init 进程探悉

前言 **2**

INIT配置文件分析 **4**

INIT的官方资料 **9**

INIT命令的手册 9
配置文件/ETC/INITTAB的手册 15

INIT详细分析 **21**

INIT PROCESS是怎么被启动的？ 21
INIT进程分析 26
init 1 的运行 *27*
主流程分析 27
辅助函数介绍 54
init 2 的运行 *68*
init 3 的运行 *74*
主流程分析 74
辅助函数 82

后记 **98**

联系 **99**

附录 **100**

环境 100
INITTAB中ACTION的注解 100
关机分析 101
关机流程介绍 *101*
Shutdown源码 *106*

前言

init是个普通的用户态进程，它是Unix系统内核初始化与用户态初始化的接合点，它是所有process的祖宗。在运行init以前是内核态初始化，该过程（内核初始化）的最后一个动作就是运行**/sbin/init**可执行文件。从init process运行开始进入Unix系统的用户态初始化。我对整个系统初始化的定义是从开机到屏幕上出现登录界面为止。这整个过程被init一分为二。当然init不单单启动了用户态的初始化，而且它在系统运行的整个期间都扮演着非常重要的角色。比如

- 在运行当中，具有 root 权限的用户可以通过再次运行 init 来切换到不同的运行级别（run level）
- init process 有认领系统中的所有孤儿进程的责任
- 当 root 权限用户想通过按 Ctrl-Alt-Del 三键来重启系统，也是由 init 最终来处理的
- 如果你想要一个 daemon 进程有这样的效果，它在整个系统运行期间一直要运行，即使它由于各种各样的原因（如在某种情况下它出错而退出了，或被某个用户 kill 掉了）停止运行了，也希望能马上被再次启动（当然不是依靠人力来手工启动），你可以在 init 运行的配置文件中加入类是与下面的一行：

  ```
  myrun::ondemand:/home/wzhou/mydaemon
  ```

  则/home/wzhou/mydaemon 这个脚本只要系统在运行，它必然也在运行。即使有人把它 kill 掉，等一会儿马上又会被 init process 启动
- 等等

而这一切都依赖于 init process。

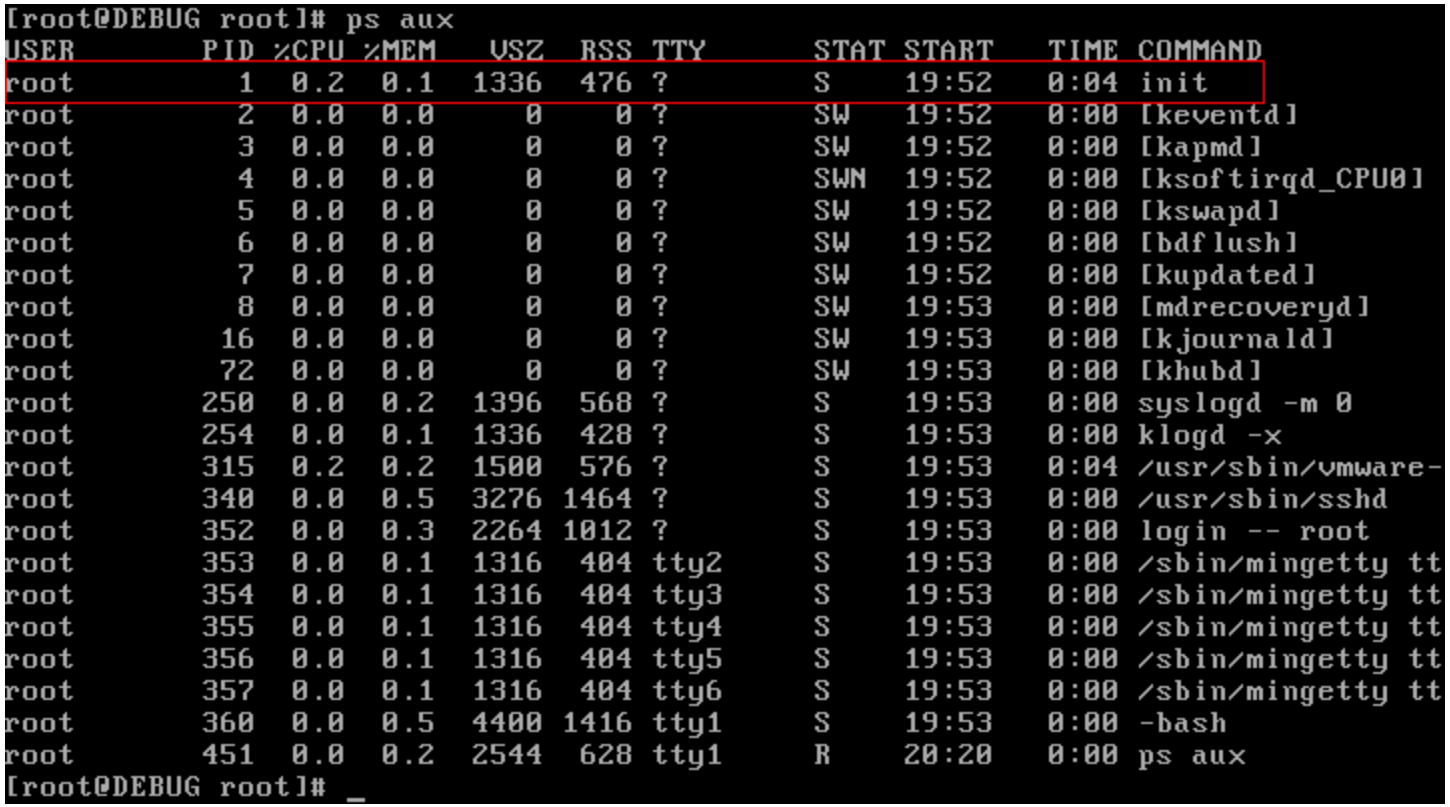

```
[root@DEBUG root]# ps aux
USER       PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1  0.2  0.1   1336   476 ?        S    19:52   0:04 init
root         2  0.0  0.0      0     0 ?        SW   19:52   0:00 [keventd]
root         3  0.0  0.0      0     0 ?        SW   19:52   0:00 [kapmd]
root         4  0.0  0.0      0     0 ?        SWN  19:52   0:00 [ksoftirqd_CPU0]
root         5  0.0  0.0      0     0 ?        SW   19:52   0:00 [kswapd]
root         6  0.0  0.0      0     0 ?        SW   19:52   0:00 [bdflush]
root         7  0.0  0.0      0     0 ?        SW   19:52   0:00 [kupdated]
root         8  0.0  0.0      0     0 ?        SW   19:53   0:00 [mdrecoveryd]
root        16  0.0  0.0      0     0 ?        SW   19:53   0:00 [kjournald]
root        72  0.0  0.0      0     0 ?        SW   19:53   0:00 [khubd]
root       250  0.0  0.2   1396   568 ?        S    19:53   0:00 syslogd -m 0
root       254  0.0  0.1   1336   428 ?        S    19:53   0:00 klogd -x
root       315  0.2  0.2   1500   576 ?        S    19:53   0:04 /usr/sbin/vmware-
root       340  0.0  0.5   3276  1464 ?        S    19:53   0:00 /usr/sbin/sshd
root       352  0.0  0.3   2264  1012 ?        S    19:53   0:00 login -- root
root       353  0.0  0.1   1316   404 tty2     S    19:53   0:00 /sbin/mingetty tt
root       354  0.0  0.1   1316   404 tty3     S    19:53   0:00 /sbin/mingetty tt
root       355  0.0  0.1   1316   404 tty4     S    19:53   0:00 /sbin/mingetty tt
root       356  0.0  0.1   1316   404 tty5     S    19:53   0:00 /sbin/mingetty tt
root       357  0.0  0.1   1316   404 tty6     S    19:53   0:00 /sbin/mingetty tt
root       360  0.0  0.5   4400  1416 tty1     S    19:53   0:00 -bash
root       451  0.0  0.2   2544   628 tty1     R    20:20   0:00 ps aux
[root@DEBUG root]# _
```

init 配置文件分析

init process的运行完全受其配置文件**/etc/inittab**的控制，这里分析一下该配置文件。

来个现实系统中的/etc/inittab 配置文件来解释一下。

```
[wzhou@dcmp10 ~]$ cat /etc/inittab
#
# inittab       This file describes how the INIT process should set up
#               the system in a certain run-level.
#
# Author:       Miquel van Smoorenburg, <miquels@drinkel.nl.mugnet.org>
#               Modified for RHS Linux by Marc Ewing and Donnie Barnes
#

# Default runlevel. The runlevels used by RHS are:
#   0 - halt (Do NOT set initdefault to this)
#   1 - Single user mode
#   2 - Multiuser, without NFS (The same as 3, if you do not have networking)
#   3 - Full multiuser mode
#   4 - unused
#   5 - X11
#   6 - reboot (Do NOT set initdefault to this)
#
id:5:initdefault:

# System initialization.
si::sysinit:/etc/rc.d/rc.sysinit

l0:0:wait:/etc/rc.d/rc 0
l1:1:wait:/etc/rc.d/rc 1
l2:2:wait:/etc/rc.d/rc 2
l3:3:wait:/etc/rc.d/rc 3
```

```
l4:4:wait:/etc/rc.d/rc 4
l5:5:wait:/etc/rc.d/rc 5
l6:6:wait:/etc/rc.d/rc 6

# Trap CTRL-ALT-DELETE
ca::ctrlaltdel:/sbin/shutdown -t3 -r now

# When our UPS tells us power has failed, assume we have a few minutes
# of power left.  Schedule a shutdown for 2 minutes from now.
# This does, of course, assume you have powerd installed and your
# UPS connected and working correctly.
pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"

# If power was restored before the shutdown kicked in, cancel it.
pr:12345:powerokwait:/sbin/shutdown -c "Power Restored; Shutdown Cancelled"


# Run gettys in standard runlevels
1:2345:respawn:/sbin/mingetty tty1
2:2345:respawn:/sbin/mingetty tty2
3:2345:respawn:/sbin/mingetty tty3
4:2345:respawn:/sbin/mingetty tty4
5:2345:respawn:/sbin/mingetty tty5
6:2345:respawn:/sbin/mingetty tty6

# Run xdm in runlevel 5
x:5:respawn:/etc/X11/prefdm -nodaemon
```

```
id:5:initdefault:
```

这一行表示系统启动后将运行在 run level 5，即 X Window 的 Full multiuser mode

```
si::sysinit:/etc/rc.d/rc.sysinit
```

sysinit 表示这是用户态系统启动，不管任何运行级别(run level)都要执行脚本/etc/rc.d/rc.sysinit
如果你要追踪操作系统内核态的初始化过程，则要从 init/main.c 中的 start_kernel()开始；而如果你想追踪操作系统用户态的启动过程，则可以从/etc/rc.d/rc.sysinit 脚本开始。

```
l0:0:wait:/etc/rc.d/rc 0      如果系统的 run level 是 0，则运行/etc/rc.d/rc 脚本，参数为 0
l1:1:wait:/etc/rc.d/rc 1      如果系统的 run level 是 1，则运行/etc/rc.d/rc 脚本，参数为 1
l2:2:wait:/etc/rc.d/rc 2      如果系统的 run level 是 2，则运行/etc/rc.d/rc 脚本，参数为 2
l3:3:wait:/etc/rc.d/rc 3      如果系统的 run level 是 3，则运行/etc/rc.d/rc 脚本，参数为 3
l4:4:wait:/etc/rc.d/rc 4      如果系统的 run level 是 4，则运行/etc/rc.d/rc 脚本，参数为 4
l5:5:wait:/etc/rc.d/rc 5      如果系统的 run level 是 5，则运行/etc/rc.d/rc 脚本，参数为 5
l6:6:wait:/etc/rc.d/rc 6      如果系统的 run level 是 6，则运行/etc/rc.d/rc 脚本，参数为 6
```

显然/etc/rc.d/rc 也是个系统初始化的很重要的脚本。上面的 wait action 表示 init process 在启动其他的动作以前，必须等待该行上的动作所代表的 process 的完成。

```
ca::ctrlaltdel:/sbin/shutdown -t3 -r now
```

这一行表示无论在什么 run level，如果 root 用户按了 Ctrl+Alt+Del 三键则运行如下命令：
/sbin/shutdown -t3 -r now
即让 init 进程监视 Ctrl+Alt+Del 事件，一旦收到，它应当运行该命令。shutdown 命令会从现在(now)开始先向系统中的所有进程发 warning，然后等待 3 秒，再杀死进程，让系统重启。

```
# When our UPS tells us power has failed, assume we have a few minutes
# of power left.  Schedule a shutdown for 2 minutes from now.
# This does, of course, assume you have powerd installed and your
# UPS connected and working correctly.
pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"
```

从上面的注释可以知道该行的动作。同样该行是不分 run level 的，只管是否发生“ powerfail ”的事件。

```
# If power was restored before the shutdown kicked in, cancel it.
pr:12345:powerokwait:/sbin/shutdown -c "Power Restored; Shutdown Cancelled"
```

在运行级别为 1,2,3,4,5 的情况下，如果发生“ powerokwait ” action，则运行命令/sbin/shutdown -c "Power Restored; Shutdown Cancelled"，即取消发出的关机指令。

```
# Run gettys in standard runlevels
1:2345:respawn:/sbin/mingetty tty1
2:2345:respawn:/sbin/mingetty tty2
3:2345:respawn:/sbin/mingetty tty3
```

```
4:2345:respawn:/sbin/mingetty tty4
5:2345:respawn:/sbin/mingetty tty5
6:2345:respawn:/sbin/mingetty tty6
```

上面的 6 行指示 init process 在 run level 是 2，3，4，5 的情况下，运行脚本/sbin/mingetty，并接受不同的参数。这里的功能是在从 tty1 到 tty6 的终端上启动字符登陆界面。

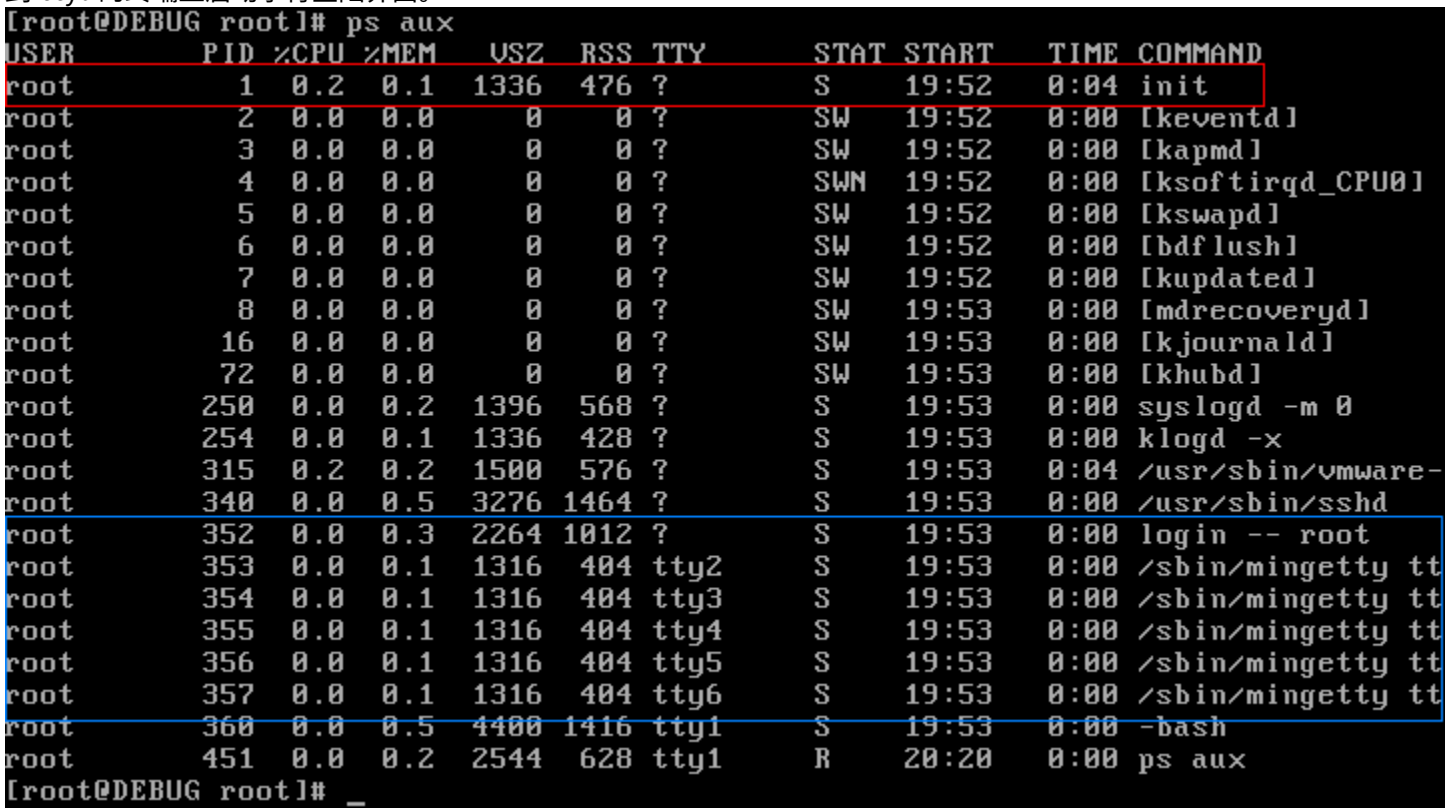

上图中用蓝框围起来的就是启动的 6 个虚拟终端。我用 root 帐号登录在 tty1,所以该终端显示“ login -- root ”,而其他 5 个虚拟终端并没有用户登录，所以还是由 mingetty 在等待着。

```
# Run xdm in runlevel 5
x:5:respawn:/etc/X11/prefdm -nodaemon
```

该行表示如果 run level 是 5，则要运行脚本/etc/X11/prefdm -nodaemon，其实就是启动 X Window,进入 GUI 界面。

上面是对 inittab 配置文件的静态的解释，下面解释 init process 依据该配置文件动态运行情况。

- init process 由“ initdefault ”知道系统将在 run level 5 下运行
- init process 首先运行“ sysinit ”标注的 action,即运行/etc/rc.d/rc.sysinit 脚本
- 运行 identifier 为“ l5 ”的动作
  l5:5:wait:/etc/rc.d/rc 5
  由于该行告诉 init process 的反映是“ wait ”，即在 init process 继续执行 inittab 配置文件中其他 action 以前，必须等待“ /etc/rc.d/rc 5 ”的结束
- 接下来执行下面的 6 个 action
  1:2345:respawn:/sbin/mingetty tty1
  2:2345:respawn:/sbin/mingetty tty2
  3:2345:respawn:/sbin/mingetty tty3
  4:2345:respawn:/sbin/mingetty tty4
  5:2345:respawn:/sbin/mingetty tty5
  6:2345:respawn:/sbin/mingetty tty6
  由于上面 6 行的 run level 告诉 init process，在 2，3，4，5 之下都要执行这里的命令“ /sbin/mingetty tty5 ”。同时这里的“ respawn ”表示如果/sbin/mingetty 所代表的 process 不运行了（无论哪种情况，是自己退出或出现问题而 crash）,init process 都有责任让他再次运行。当启动 Linux 后我们通过 Alt-F1,…,Alt-F6 可以切换到相应的终端，就是这几行运行的缘故。另外，当你登录到某个终端，比如 tty1,然后在命令行上输入 exit,在该终端上又会出现登录界面，这就是 init process 在响应“ respawn ”动作。当你输入 exit 时，/sbin/mingetty 代表的 process 退出，被 init process 监控到，马上在该终端上又运行“ /sbin/mingetty tty1 ”，从而在退出的 tty1 上再次出现登录界面。
- 最后运行的是/etc/X11/prefdm -nodaemon，即启动 X Window 登录。
- 在配置文件中的下面的配置行并不会执行，但会被 init process 纪录状态。只有当系统出现对应的情况时，才会运行。
  ca::ctrlaltdel:/sbin/shutdown -t3 -r now
  pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"
  pr:12345:powerokwait:/sbin/shutdown -c "Power Restored; Shutdown Cancelled"

  比如当 root 用户按了 Ctrl-Alt-Del 键以后，init process 将执行如下命令行“ /sbin/shutdown -t3 -r now ”进行关机；而当 UPS 报告电源出现故障，马上要断电时，就执行“ /sbin/shutdown -f -h +2 "Power Failure; System Shutting Down" ”；当 UPS 报告从电源故障恢复以后，执行“ /sbin/shutdown -c "Power Restored; Shutdown Cancelled" ”。那么 init process 是怎么感知这些消息的呢？即该进程怎么知道 root 权限用户按下了 Ctrl-Alt-Del 键，UPS 报告电源出现故障及电源恢复呢？都是通过 Unix 特有的 signal 机制。对 init process 来说，它只要正确处理对应的 signal 就好。

init 的官方资料

init 的作者亲手写了与 init 相关的手册，即 man init 与 man inittab。仔细看看，对理解 init process 有很大帮助。

init 命令的手册

```
NAME
       init, telinit - process control initialization

SYNOPSIS
       /sbin/init [ -a ] [ -s ] [ -b ] [ -z xxx ] [ 0123456Ss ]
       /sbin/telinit [ -t sec ] [ 0123456sSQqabcUu ]

DESCRIPTION
   Init
       Init  is the parent of all processes.  Its primary role is to create processes
       from a script stored in the file /etc/inittab  (see  inittab(5)).   This  file
       usually  has  entries which cause init to spawn gettys on each line that users
       can log in.  It also controls autonomous processes required by any  particular
       system.

RUNLEVELS
       A  runlevel  is  a  software  configuration  of the system which allows only a
       selected group of processes to exist.  The processes spawned by init for  each
       of  these  runlevels are defined in the /etc/inittab file.  Init can be in one
       of eight runlevels: 0â€“6 and S or s.  The runlevel is changed by having a priv-
       ileged  user  run telinit, which sends appropriate signals to init, telling it
       which runlevel to change to.
```

Runlevels 0, 1, and 6 are reserved. Runlevel 0 is used to halt the system, runlevel 6 is used to reboot the system, and runlevel 1 is used to get the system down into single user mode. Runlevel S is not really meant to be used directly, but more for the scripts that are executed when entering runlevel 1. For more information on this, see the manpages for shutdown(8) and inittab(5).

Runlevels 7-9 are also valid, though not really documented. This is because "traditional" Unix variants donâ€™t use them. In case youâ€™re curious, runlevels S and s are in fact the same. Internally they are aliases for the same runlevel.

BOOTING

After init is invoked as the last step of the kernel boot sequence, it looks for the file /etc/inittab to see if there is an entry of the type initdefault (see inittab(5)). The initdefault entry determines the initial runlevel of the system. If there is no such entry (or no /etc/inittab at all), a runlevel must be entered at the system console.

Runlevel S or s bring the system to single user mode and do not require an /etc/inittab file. In single user mode, a root shell is opened on /dev/console.

When entering single user mode, init initializes the consoles stty settings to sane values. Clocal mode is set. Hardware speed and handshaking are not changed.

When entering a multi-user mode for the first time, init performs the boot and bootwait entries to allow file systems to be mounted before users can log in. Then all entries matching the runlevel are processed.

When starting a new process, init first checks whether the file /etc/initscript exists. If it does, it uses this script to start the process.

Each time a child terminates, init records the fact and the reason it died in /var/run/utmp and /var/log/wtmp, provided that these files exist.

CHANGING RUNLEVELS

After it has spawned all of the processes specified, init waits for one of its descendant processes to die, a powerfail signal, or until it is signaled by telinit to change the systemâ€™s runlevel. When one of the above three conditions occurs, it re-examines the /etc/inittab file. New entries can be added to this file at any time. However, init still waits for one of the above three conditions to occur. To provide for an instantaneous response, the telinit Q or q command can wake up init to re-examine the /etc/inittab file.

If init is not in single user mode and receives a powerfail signal (SIGPWR), it reads the file /etc/powerstatus. It then starts a command based on the contents of this file:

F(AIL) Power is failing, UPS is providing the power. Execute the powerwait and powerfail entries.

O(K) The power has been restored, execute the powerokwait entries.

L(OW) The power is failing and the UPS has a low battery. Execute the powerfailnow entries.

If /etc/powerstatus doesnâ€™t exist or contains anything else then the letters F, O or L, init will behave as if it has read the letter F.

Usage of SIGPWR and /etc/powerstatus is discouraged. Someone wanting to interact with init should use the /dev/initctl control channel - see the source code of the sysvinit package for more documentation about this.

When init is requested to change the runlevel, it sends the warning signal SIGTERM to all processes that are undefined in the new runlevel. It then waits 5 seconds before forcibly terminating these processes via the SIGKILL signal. Note that init assumes that all these processes (and their descendants) remain in the same process group which init originally created for them. If any process changes its process group affiliation it will not receive these signals. Such processes need to be terminated separately.

TELINIT

/sbin/telinit is linked to /sbin/init. It takes a one-character argument and signals init to perform the appropriate action. The following arguments serve as directives to telinit:

0,1,2,3,4,5 or 6
tell init to switch to the specified run level.

a,b,c tell init to process only those /etc/inittab file entries having run-level a,b or c.

Q or q tell init to re-examine the /etc/inittab file.

S or s tell init to switch to single user mode.

U or u tell init to re-execute itself (preserving the state). No re-examining of /etc/inittab file happens. Run level should be one of Ss12345, otherwise request would be silently ignored.

telinit can also tell init how long it should wait between sending processes the SIGTERM and SIGKILL signals. The default is 5 seconds, but this can be changed with the -t sec option.

telinit can be invoked only by users with appropriate privileges.

The init binary checks if it is init or telinit by looking at its process id; the real initâ€™s process id is always 1. From this it follows that instead of calling telinit one can also just use init instead as a shortcut.

ENVIRONMENT

Init sets the following environment variables for all its children:

PATH /usr/local/sbin:/sbin:/bin:/usr/sbin:/usr/bin

INIT_VERSION
As the name says. Useful to determine if a script runs directly from init.

RUNLEVEL
The current system runlevel.

PREVLEVEL
The previous runlevel (useful after a runlevel switch).

CONSOLE
The system console. This is really inherited from the kernel; however if it is not set init will set it to /dev/console by default.

BOOTFLAGS

It is possible to pass a number of flags to init from the boot monitor (eg. LILO). Init accepts the following flags:

-s, S, single
Single user mode boot. In this mode /etc/inittab is examined and the bootup rc scripts are usually run before the single user mode shell is started.

1-5 Runlevel to boot into.

-b, emergency
Boot directly into a single user shell without running any other startup scripts.

-a, auto
The LILO boot loader adds the word "auto" to the command line if it booted the kernel with the default command line (without user intervention). If this is found init sets the "AUTOBOOT" environment variable to "yes". Note that you cannot use this for any security measures - of course the user could specify "auto" or -a on the command line manually.

-z xxx
The argument to -z is ignored. You can use this to expand the command line a bit, so that it takes some more space on the stack. Init can then manipulate the command line so that ps(1) shows the current runlevel.

INTERFACE

Init listens on a fifo in /dev, /dev/initctl, for messages. Telinit uses this to communicate with init. The interface is not very well documented or finished. Those interested should study the initreq.h file in the src/ subdirectory of the init source code tar archive.

SIGNALS

Init reacts to several signals:

SIGHUP

Has the same effect as telinit q.

SIGUSR1

On receipt of this signals, init closes and re-opens its control fifo, /dev/initctl. Useful for bootscripts when /dev is remounted.

SIGINT

Normally the kernel sends this signal to init when CTRL-ALT-DEL is pressed. It activates the ctrlaltdel action.

SIGWINCH

The kernel sends this signal when the KeyboardSignal key is hit. It activates the kbrequest action.

CONFORMING TO

Init is compatible with the System V init. It works closely together with the scripts in the directories /etc/init.d and /etc/rc{runlevel}.d. If your system uses this convention, there should be a README file in the directory /etc/init.d explaining how these scripts work.

FILES

/etc/inittab
/etc/initscript
/dev/console
/var/run/utmp
/var/log/wtmp
/dev/initctl

```
WARNINGS
       Init  assumes  that  processes and descendants of processes remain in the same
       process group which was originally created for them.  If the processes  change
       their  group, init canâ€™t kill them and you may end up with two processes read-
       ing from one terminal line.

DIAGNOSTICS
       If init finds that it is continuously respawning an entry more than  10  times
       in  2  minutes,  it  will assume that there is an error in the command string,
       generate an error message on the system console, and refuse  to  respawn  this
       entry  until  either 5 minutes has elapsed or it receives a signal.  This pre-
       vents it from eating up system resources when someone  makes  a  typographical
       error in the /etc/inittab file or the program for the entry is removed.

AUTHOR
       Miquel  van  Smoorenburg  (miquels@cistron.nl), initial manual page by Michael
       Haardt (u31b3hs@pool.informatik.rwth-aachen.de).

SEE ALSO
       getty(1), login(1),  sh(1),  runlevel(8),  shutdown(8),  kill(1),  inittab(5),
       initscript(5), utmp(5)

                                    18 April 2003                              INIT(8)
```

配置文件/etc/inittab 的手册

```
INITTAB(5)                      Linux System Administrator's Manual                      INITTAB(5)

NAME
       inittab - format of the inittab file used by the sysv-compatible init process

DESCRIPTION
       The  inittab  file  describes which processes are started at bootup and during normal operation
       (e.g. /etc/init.d/boot, /etc/init.d/rc, gettys...).  Init(8) distinguishes multiple  runlevels,
```

each of which can have its own set of processes that are started. Valid runlevels are 0-6 plus A, B, and C for ondemand entries. An entry in the inittab file has the following format:

```
id:runlevels:action:process
```

Lines beginning with '#' are ignored.

id
: is a unique sequence of 1-4 characters which identifies an entry in inittab (for versions of sysvinit compiled with the old libc5 (< 5.2.18) or a.out libraries the limit is 2 characters).

  Note: traditionally, for getty and other login processes, the value of the id field is kept the same as the suffix of the corresponding tty, e.g. 1 for tty1. Some ancient login accounting programs might expect this, though I can't think of any.

runlevels
: lists the runlevels for which the specified action should be taken.

action describes which action should be taken.

process
: specifies the process to be executed. If the process field starts with a '+' character, init will not do utmp and wtmp accounting for that process. This is needed for gettys that insist on doing their own utmp/wtmp housekeeping. This is also a historic bug.

The runlevels field may contain multiple characters for different runlevels. For example, 123 specifies that the process should be started in runlevels 1, 2, and 3. The runlevels for ondemand entries may contain an A, B, or C. The runlevels field of sysinit, boot, and bootwait entries are ignored.

When the system runlevel is changed, any running processes that are not specified for the new runlevel are killed, first with SIGTERM, then with SIGKILL.

Valid actions for the action field are:

respawn
: The process will be restarted whenever it terminates (e.g. getty).

wait The process will be started once when the specified runlevel is entered and init will wait for its termination.

once The process will be executed once when the specified runlevel is entered.

boot The process will be executed during system boot. The runlevels field is ignored.

bootwait
The process will be executed during system boot, while init waits for its termination (e.g. /etc/rc). The runlevels field is ignored.

off This does nothing.

ondemand[1]
A process marked with an ondemand runlevel will be executed whenever the specified ondemand runlevel is called. However, no runlevel change will occur (ondemand runlevels are 'a', 'b', and 'c').

initdefault
An initdefault entry specifies the runlevel which should be entered after system boot. If none exists, init will ask for a runlevel on the console. The process field is ignored.

sysinit
The process will be executed during system boot. It will be executed before any boot or bootwait entries. The runlevels field is ignored.

powerwait
The process will be executed when the power goes down. Init is usually informed about this by a process talking to a UPS connected to the computer. Init will wait for the process to finish before continuing.

powerfail
As for powerwait, except that init does not wait for the process's completion.

[1] Ondemand与respawn的区别是其与运行级别（run level）无关

powerokwait
This process will be executed as soon as init is informormed that the power has been restored.

powerfailnow
This process will be executed when init is told that the battery of the external UPS is almost empty and the power is failing (provided that the external UPS and the monitoring process are able to detect this condition).

ctrlaltdel
The process will be executed when init receives the SIGINT signal. This means that someone on the system console has pressed the CTRL-ALT-DEL key combination. Typically one wants to execute some sort of shutdown either to get into single-user level or to reboot the machine.

kbrequest
The process will be executed when init receives a signal from the keyboard handler that a special key combination was pressed on the console keyboard.

The documentation for this function is not complete yet; more documentation can be found in the kbd-x.xx packages (most recent was kbd-0.94 at the time of this writing). Basically you want to map some keyboard combination to the "KeyboardSignal" action. For example, to map Alt-Uparrow for this purpose use the following in your keymaps file:

```
alt keycode 103 = KeyboardSignal
```

EXAMPLES
This is an example of a inittab which resembles the old Linux inittab:

```
# inittab for linux
id:1:initdefault:
rc::bootwait:/etc/rc
1:1:respawn:/etc/getty 9600 tty1
2:1:respawn:/etc/getty 9600 tty2
3:1:respawn:/etc/getty 9600 tty3
4:1:respawn:/etc/getty 9600 tty4
```

```
This inittab file executes /etc/rc during boot and starts gettys on tty1-tty4.

A more elaborate inittab with different runlevels (see the comments inside):

        # Level to run in
        id:2:initdefault:

        # Boot-time system configuration/initialization script.
        si::sysinit:/etc/init.d/rcS

        # What to do in single-user mode.
        ~:S:wait:/sbin/sulogin

        # /etc/init.d executes the S and K scripts upon change
        # of runlevel.
        #
        # Runlevel 0 is halt.
        # Runlevel 1 is single-user.
        # Runlevels 2-5 are multi-user.
        # Runlevel 6 is reboot.

        l0:0:wait:/etc/init.d/rc 0
        l1:1:wait:/etc/init.d/rc 1
        l2:2:wait:/etc/init.d/rc 2
        l3:3:wait:/etc/init.d/rc 3
        l4:4:wait:/etc/init.d/rc 4
        l5:5:wait:/etc/init.d/rc 5
        l6:6:wait:/etc/init.d/rc 6

        # What to do at the "3 finger salute".
        ca::ctrlaltdel:/sbin/shutdown -t1 -h now

        # Runlevel 2,3: getty on virtual consoles
        # Runlevel   3: getty on terminal (ttyS0) and modem (ttyS1)
        1:23:respawn:/sbin/getty tty1 VC linux
        2:23:respawn:/sbin/getty tty2 VC linux
```

```
              3:23:respawn:/sbin/getty tty3 VC linux
              4:23:respawn:/sbin/getty tty4 VC linux
              S0:3:respawn:/sbin/getty -L 9600 ttyS0 vt320
              S1:3:respawn:/sbin/mgetty -x0 -D ttyS1

FILES
       /etc/inittab

AUTHOR
       Init  was written by Miquel van Smoorenburg (miquels@cistron.nl).  This manual page was written
       by Sebastian Lederer (lederer@francium.informatik.uni-bonn.de) and modified by  Michael  Haardt
       (u31b3hs@pool.informatik.rwth-aachen.de).

SEE ALSO
       init(8), telinit(8)
```

init 详细分析

init process 是怎么被启动的？

init process 是 Linux 系统的第一个用户态进程，那自然没有父亲。它是由 Linux 内核直接启动的。

start_kernel()是内核的汇编与 C 语言的交接点，在该函数以前，内核的代码都是用汇编写的，完成一些最基本的初始化与环境设置工作，比如内核代码载入内存并解压缩（现在的内核一般都经过压缩），CPU 的最基本初始化，为 C 代码的运行设置环境（C 代码的运行是有一定环境要求的，比如 stack 的设置等）。这里一个不太确切的比喻是 start_kernel()就像是 C 代码中的 main()。我们知道对应用程序员而言，main()是他的入口，但实际上程序的入口是被包在了 C 库中，在链接阶段，linker 会把它链接入你的程序中。而它的任务中有一项就是为 main()准备运行环境。main()中的 argc，argv 等都不是平白无故来的，都是在调用 main()以前的代码做的准备。

在 start_kernel()中 Linux 将完成整个系统的内核初始化。内核初始化的最后一步就是启动 init 进程这个所有进程的祖先。

Linux-2.6.20/init/main.c

```
483     asmlinkage void __init start_kernel(void)       该函数是 Linux 内核的入口，其前面的代码都是用汇编编写
484     {
485           char * command_line;
486           extern struct kernel_param __start___param[], __stop___param[];
487
488           smp_setup_processor_id();
489
490           /*
491            * Need to run as early as possible, to initialize the
492            * lockdep hash:
493            */
494           unwind_init();
```

```
495             lockdep_init();
496
497             local_irq_disable();
498             early_boot_irqs_off();
499             early_init_irq_lock_class();
500
501     /*
502      * Interrupts are still disabled. Do necessary setups, then
503      * enable them
504      */
505             lock_kernel();
506             boot_cpu_init();
507             page_address_init();
508             printk(KERN_NOTICE);
509             printk(linux_banner);

                ∘ ∘ ∘ ∘ ∘ ∘

611             cpuset_init();
612             taskstats_init_early();
613             delayacct_init();
614
615             check_bugs();
616
617             acpi_early_init(); /* before LAPIC and SMP init */
618
619             /* Do the rest non-__init'ed, we're now alive */
620             rest_init();                    这是 Linux 内核初始化的尾声
621     }
```

```
416     static void noinline rest_init(void)
417             __releases(kernel_lock)
418     {
419             kernel_thread(init, NULL, CLONE_FS | CLONE_SIGHAND);   创建一个内核线程，实际上就是内核进程，
420             numa_default_policy();                                 Linux 内核是不支持类似 Windows NT 一样的
421             unlock_kernel();                                       线程概念的。Linux 本质上只支持进程。这里的
```

```
422                                                          init 只是一个函数，不要与 init process 搞
423             /*                                           混淆了。该函数见下面。
424              * The boot idle thread must execute schedule()
425              * at least one to get things moving:
426              */
427             preempt_enable_no_resched();
428             schedule();
429             preempt_disable();
430
431             /* Call into cpu_idle with preempt disabled */
432             cpu_idle();
433     }
```

```
716     static int init(void * unused)      内核创建的内核线程运行本函数，在本函数的结尾就是启动 init process
717     {
718             lock_kernel();
719             /*
720              * init can run on any cpu.
721              */
722             set_cpus_allowed(current, CPU_MASK_ALL);
723             /*
724              * Tell the world that we're going to be the grim
725              * reaper of innocent orphaned children.
726              *
727              * We don't want people to have to make incorrect
728              * assumptions about where in the task array this
729              * can be found.
730              */
731             init_pid_ns.child_reaper = current;
732
733             cad_pid = task_pid(current);
734
735             smp_prepare_cpus(max_cpus);
736
737             do_pre_smp_initcalls();
738
```

```
739             smp_init();
740             sched_init_smp();
741
742             cpuset_init_smp();
743
744             do_basic_setup();
745
746             /*
747              * check if there is an early userspace init.  If yes, let it do all
748              * the work
749              */
750
751             if (!ramdisk_execute_command)
752                   ramdisk_execute_command = "/init";
753
754             if (sys_access((const char __user *) ramdisk_execute_command, 0) != 0) {
755                   ramdisk_execute_command = NULL;
756                   prepare_namespace();
757             }
758
759             /*
760              * Ok, we have completed the initial bootup, and
761              * we're essentially up and running. Get rid of the
762              * initmem segments and start the user-mode stuff..
763              */
764             free_initmem();
765             unlock_kernel();
766             mark_rodata_ro();
767             system_state = SYSTEM_RUNNING;
768             numa_default_policy();
769
770             if (sys_open((const char __user *) "/dev/console", O_RDWR, 0) < 0)
771                   printk(KERN_WARNING "Warning: unable to open an initial console.\n");
772
773             (void) sys_dup(0);
774             (void) sys_dup(0);
775
```

```
776             if (ramdisk_execute_command) {
777                   run_init_process(ramdisk_execute_command);
778                   printk(KERN_WARNING "Failed to execute %s\n",
779                                 ramdisk_execute_command);
780             }
781
782             /*
783              * We try each of these until one succeeds.
784              *
785              * The Bourne shell can be used instead of init if we are
786              * trying to recover a really broken machine.
787              */
788             if (execute_command) {
789                   run_init_process(execute_command);
790                   printk(KERN_WARNING "Failed to execute %s.  Attempting "
791                                       "defaults...\n", execute_command);
792             }
793             run_init_process("/sbin/init");
794             run_init_process("/etc/init");
795             run_init_process("/bin/init");
796             run_init_process("/bin/sh");
797
798             panic("No init found.  Try passing init= option to kernel.");
799       }
```

run_init_process()实际上是通过嵌入汇编构建一个类似用户态代码一样的 sys_execve()调用，其参数就是要执行的可执行文件名，也就是这里的 init process 在磁盘上的文件。

这里的 run_init_process 就是通过 execve()来运行 init 程序。这里首先运行“ /sbin/init ”，如果失败再运行“ /etc/init ”,然后是“ /bin/init ”,然后是“ /bin/sh ”（也就是说，init 可执行文件可以放在上面代码中寻找的 4 个目录中都可以）,如果都失败，则可以通过在系统启动时在添加的启动参数来指定 init，比如 init=/home/wzhou/init。这里是内核初始化结束并开始用户态初始化的阴阳界。

```
710       static void run_init_process(char *init_filename)
711       {
712             argv_init[0] = init_filename;
713             kernel_execve(init_filename, argv_init, envp_init);
714       }
```

Linux-2.6.20/arch/i386/kernel/sys_i386.c

```
254     /*
255      * Do a system call from kernel instead of calling sys_execve so we
256      * end up with proper pt_regs.
257      */    构造 sys_execve 系统调用
258     int kernel_execve(const char *filename, char *const argv[], char *const envp[])
259     {
260             long __res;
261             asm volatile ("push %%ebx ; movl %2,%%ebx ; int $0x80 ; pop %%ebx"
262             : "=a" (__res)
263             : "0" (__NR_execve),"ri" (filename),"c" (argv), "d" (envp) : "memory");
264             return __res;
265     }
```

init 进程分析

整个 init 的代码比较难读，倒不是真的 init process 要完成的工作有多么复杂，在我看来，这复杂大半的原因是设计者自找的。

init 的执行常规分成三种：

1. 在系统启动阶段，操作系统内核部分初始化阶段的结尾，将运行 init 这个第一个用户态的程序（它将作为所有用户态进程的共同祖先），它将依据/etc/inittab 配置文件来对系统进行用户态的初始化。
2. 在系统运行当中 root 用户可以运行 init 命令把系统切换到不同的运行级别（run level）。比如当前运行级别是 3（Console 界面的 Full multiuser mode），而 root 想维护系统，他可以运行如下命令：
   # init 1　　切换到 Single user mode,即单用户模式，有点像 Windows 下的安全模式
   用户启动的 init 命令并不真正运行 run level 切换的工作，只是通过 pipe(管道)把命令打包成 request，然后传递给作为 daemon 进程运行的 init。
3. 在系统起来以后，init 作为一个 daemon 进程运行，一是监控/etc/inittab 配置文件中的相关命令的执行，二就是通过 pipe(管道)接受 2 中发来的切换 run level 的 request(请求)并处理之。

设计者把上面的功能合三为一，把整个逻辑完全搅和在一起，造成代码的难读。我不敢怀疑代码作者的水平，我只是想这可能是 Linux 下有些开发者的特点。就像 Linux 之父非常坚定地拒绝把 kernel 级的 debugger 支持引入其一手遮天的官方内核，虽然他有他的理由，但我想很多 Linux 下的内核黑客并不一定认同他的理由，只能无奈的接受他的这个“特点”。

我把 init process 的三种状态分别称为“init 1”，“init 2”，“init 3”，分别对应上面列举的三种状态。如果把这三种状态的 init 混在一块儿说的话，很容易搞成一团乱麻。我在下面把 init 按三种角色来说明，虽然实际上只有一个可执行文件，进程常规情况下是一个，当通过 init 切换 run level 时会是两个，绝不会是三个。

init 1 的运行

主流程分析

init 1 是由内核启动的，不带任何命令行参数，即直接执行/sbin/init。

```
2594  /*
2595   * Main entry for init and telinit.
2596   */
2597  int main(int argc, char **argv)          这时 argc 为 1，argv[0] = “/sbin/init”
2598  {
2599        char              *p;
2600        int               f;
2601        int               isinit;
2602
2603        /* Get my own name */
2604        if ((p = strrchr(argv[0], '/')) != NULL)       argv[0] = /sbin/init
2605              p++;                                     则 p 指向 init
2606        else
2607              p = argv[0];
2608        umask(022);
2609
2610        /* Quick check */
2611        if (geteuid() != 0) {         检查是否拥有 root 权限，运行 init 必须拥有该权限，内核当然拥有 root 权限
```

```
2612                fprintf(stderr, "%s: must be superuser.\n", p);
2613                exit(1);
2614        }
2615
2616        /*
2617         *    Is this telinit or init ?
2618         */
2619        isinit = (getpid() == 1);           内核启动的 init process 的 PID 为 1，满足，即 isinit = true

2620        for (f = 1; f < argc; f++) {        init 1 没有命令行参数，所以不尽如循环
2621                if (!strcmp(argv[f], "-i") || !strcmp(argv[f], "--init"))
2622                        isinit = 1;
2623                        break;
2624        }
2625        if (!isinit) exit(telinit(p, argc, argv));      init 1 不执行
2626
2627        /*
2628         *    Check for re-exec
2629         */
2630        if (check_pipe(STATE_PIPE)) {               检查 init 1 与 init 3 之间沟通的 pipe 是否建立，这时是内核启动
2631                                                    的 init，当然还白废待新，什么都没有呢。所以这里 check_pipe()
2632                receive_state(STATE_PIPE);          返回 0，init 0 不会进入该 if 分支，跳到 L2646 执行
2633
2634                myname = istrdup(argv[0]);
2635                argv0 = argv[0];
2636                maxproclen = 0;
2637                for (f = 0; f < argc; f++)
2638                        maxproclen += strlen(argv[f]) + 1;
2639                reload = 1;
2640                setproctitle("init [%c]",runlevel);
2641
2642                init_main();
2643        }
2644
2645        /* Check command line arguments */      由于启动 init 1 时没有带任何命令行参数，所以不会进入下面的循环
2646        maxproclen = strlen(argv[0]) + 1;       直接跳到 L2666 执行
```

```
2647        for(f = 1; f < argc; f++) {
2648              if (!strcmp(argv[f], "single") || !strcmp(argv[f], "-s"))
2649                    dfl_level = 'S';
2650              else if (!strcmp(argv[f], "-a") || !strcmp(argv[f], "auto"))
2651                    putenv("AUTOBOOT=YES");
2652              else if (!strcmp(argv[f], "-b") || !strcmp(argv[f],"emergency"))
2653                    emerg_shell = 1;
2654              else if (!strcmp(argv[f], "-z")) {
2655                    /* Ignore -z xxx */
2656                    if (argv[f + 1]) f++;
2657              } else if (strchr("0123456789sS", argv[f][0])
2658                    && strlen(argv[f]) == 1)
2659                    dfl_level = argv[f][0];
2660              /* "init u" in the very beginning makes no sense */
2661              if (dfl_level == 's') dfl_level = 'S';
2662              maxproclen += strlen(argv[f]) + 1;
2663        }
2664
2665        /* Start booting. */
2666        argv0 = argv[0];                    到此，init 1 的 argv0 = /sbin/init
2667        argv[1] = NULL;
2668        setproctitle("init boot");
2669        init_main(dfl_level);               init 1 调用 init_main(0)，dfl_level 被静态的初始化为 0
2670
2671        /*NOTREACHED*/
2672        return 0;
2673  }
```

OK，init 1 进入主函数 init_main()。

```
2340  /*
2341   *    The main loop
2342   */
2343  int init_main()
2344  {
2345    CHILD                 *ch;
2346    struct sigaction      sa;
2347    sigset_t              sgt;
```

```
2348    pid_t                   rc;
2349    int             f, st;
2350
2351    if (!reload) {              init 1 不会修改 reload 值，所以还是 0，所以会进入这里的 if 分支
2352
2353  #if INITDEBUG                 debug init 用，忽略，debug init 1 可真有点技巧。因为/sbin/init 虽然是普通的用户
2354        /*                      程序，但别忘了它运行时的时机，debugger 根本还没机会介入呢
2355         * Fork so we can debug the init process.
2356         */
2357        if ((f = fork()) > 0) {
2358              static const char killmsg[] = "PRNT: init killed.\r\n";
2359              pid_t rc;
2360
2361              while((rc = wait(&st)) != f)
2362                    if (rc < 0 && errno == ECHILD)
2363                          break;
2364              write(1, killmsg, sizeof(killmsg) - 1);
2365              while(1) pause();
2366        }
2367  #endif
2368
2369  #ifdef __linux__              因为该 init 的代码,FreeBSD 也会用到，所以用该 Macro 来表示是为 Linux 编译的
2370        /*
2371         *    Tell the kernel to send us SIGINT when CTRL-ALT-DEL
2372         *    is pressed, and that we want to handle keyboard signals.
2373         */
2374        init_reboot(BMAGIC_SOFT);           调用 reboot(BMAGIC_SOFT)，使得当按下 CTRL-ALT-DEL 后，，将向 init
                                              process发SIGINT signal，见man 2 reboot

2375        if ((f = open(VT_MASTER, O_RDWR | O_NOCTTY)) >= 0) {
2376              (void) ioctl(f, KDSIGACCEPT, SIGWINCH);
2377              close(f);
2378        } else
2379              (void) ioctl(0, KDSIGACCEPT, SIGWINCH);
2380  #endif
2381
```

```
2382        /*
2383         *    Ignore all signals.          把当前所有 signal handler 都设为 ignore，即不处理该 signals，因为下
2384         */                                面要重设这些 signals 的 handler
2385        for(f = 1; f <= NSIG; f++)
2386              SETSIG(sa, f, SIG_IGN, SA_RESTART);
2387    }
2388
```

背景介绍

对需要特殊处理的 signal 进行设置。
SIGALRM 为超时信号，即设置好时钟，当时钟到时后发该信号
SIGHUP 为连接断开信号，比如你通过 telnet 远程登陆到某台 Linux 机器上，启动 top 程序，然后从 telnet 退出，这是你刚才启动的 top 程序会收到该信号。
SIGINT 为中断操作信号，当用户按 Ctrl-C 时，前台进程组会收到该信号,系统将会把 CTRL-ALT-DEL 转换成该 signal
SIGCHLD 为当进程被终止或停止时会发该信号给其父进程
SIGPWR 为当电源失效，UPS 开始工作时，系统会发该信号给 init 进程
SIGWINCH（WINdow CHange）不太了了
SIGUSR1 为用户定义信号
SIGSTOP 停止信号
SIGTSTP 交互停止信号，当用户在终端上按 Ctrl-Z 后，当前进程会挂起
SIGCONT 为与 SIGSTOP 相对的 continue 信号
SIGSEGV 段违例，一般指访问了非法内存

当 init process 收到 SIGALRM
SIGHUP
SIGINT
SIGPWR
SIGWINCH
SIGUSR1 时会执行 signal_handler（），signal_handler（）只是把收到的 signal 记录在全局变量 got_signals 中了事。

当 init process 收到 SIGCHLD，会执行 chld_handler（），即当 init process 的子进程死亡时，会执行 chld_handler（），取得该死亡儿子的退出码。

当 init process 收到 SIGSTOP，SIGTSTP，会执行 stop_handler ()
当 init process 收到 SIGCONT，会执行 cont_handler ()
当 init process 收到 SIGSEGV，会执行 segv_handler ()，即 init process 访问非法内存后的处理，只是 init process 睡眠３０秒，然后接着干。一般 process 如果访问非法内存，都会死掉。

```
2389    SETSIG(sa, SIGALRM,  signal_handler, 0);                    重设 init process 关心的 signal handler
2390    SETSIG(sa, SIGHUP,   signal_handler, 0);
2391    SETSIG(sa, SIGINT,   signal_handler, 0);
2392    SETSIG(sa, SIGCHLD,  chld_handler, SA_RESTART);
2393    SETSIG(sa, SIGPWR,   signal_handler, 0);
2394    SETSIG(sa, SIGWINCH, signal_handler, 0);
2395    SETSIG(sa, SIGUSR1,  signal_handler, 0);
2396    SETSIG(sa, SIGSTOP,  stop_handler, SA_RESTART);
2397    SETSIG(sa, SIGTSTP,  stop_handler, SA_RESTART);
2398    SETSIG(sa, SIGCONT,  cont_handler, SA_RESTART);
2399    SETSIG(sa, SIGSEGV,  (void (*)(int))segv_handler, SA_RESTART);
```

让我们看看这里注册的几个 signal handler 到底干了点什么？
响应 SIGALRM，SIGHUP，SIGINT，SIGPWR，SIGWINCH，SIGUSR1 的是 signal_handler。

```
543   /*
544    *    We got a signal (HUP PWR WINCH ALRM INT)
545    */
546   void signal_handler(int sig)
547   {
548         ADDSET(got_signals, sig);          对于 HUP PWR WINCH ALRM INT  signal 只是记录一下了事，具体处理在
549   }                                        init_main()中的 process_signals()
```

got_signals 是一个全局变量。

```
106   sig_atomic_t got_signals;     /* Set if we received a signal. */
```

而 ADDSET () 只是用来把发生的 signal 纪录到该全局变量中。

```
#define ADDSET(set, val)   ((set) |=  (1 << (val)))
```

对这些纪录在 got_signals 变量中的 signal 的处理在函数 process_signals()中。

```
2238  void process_signals()              处理 init process 受到的 SIGALRM，SIGHUP，SIGINT，SIGPWR，SIGWINCH，
2239  {                                   SIGUSR1 signal
```

```
2240    CHILD           *ch;
2241    int       pwrstat;
2242    int       oldlevel;
2243    int       fd;
2244    char            c;
2245
2246    if (ISMEMBER(got_signals, SIGPWR)) {    收到过 SIGPWR signal，即 UPS 报告电源 fail
2247        INITDBG(L_VB, "got SIGPWR");
2248        /* See _what_ kind of SIGPWR this is. */
2249        pwrstat = 0;
2250        if ((fd = open(PWRSTAT, O_RDONLY)) >= 0) {      打开/etc/powerstatus 文件，如果该文件存在，则可能是如
2251            c = 0;                                      下三个字符中的一个：“F”，“O”，“L”，应该是表示收到
2252            read(fd, &c, 1);                            SIGPWR  signal 的原因吧。F 表示 fail，O 表示 OK，L 表示 low

2253            pwrstat = c;                        把电源 fail 的原因纪录在变量 pwrstat 中
2254            close(fd);
2255            unlink(PWRSTAT);                    删除/etc/powerstatus 文件
2256        }
2257        do_power_fail(pwrstat);       根据 powerfail 的原因来对 family 链表中的 action 进行处理，而该连表中的
                                          action 是完全按照配置文件 inittab 中的“powerfail”配置来实行的，见下面的

            pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"
            这里 init process 就会执行/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"

2258        DELSET(got_signals, SIGPWR);  从 got_signals 全局变量中删去 SIGPWR 标志
2259    }
2260
2261    if (ISMEMBER(got_signals, SIGINT)) {    收到过 SIGINT  signal
2262        INITDBG(L_VB, "got SIGINT");
2263        /* Tell ctrlaltdel entry to start up */
2264        for(ch = family; ch; ch = ch->next)
2265            if (ch->action == CTRLALTDEL)
2266                ch->flags &= ~XECUTED;              允许 Ctrl-Alt-Del handler 的运行
2267        DELSET(got_signals, SIGINT);
2268    }
```

```
2269
2270    if (ISMEMBER(got_signals, SIGWINCH)) {
2271        INITDBG(L_VB, "got SIGWINCH");
2272        /* Tell kbrequest entry to start up */
2273        for(ch = family; ch; ch = ch->next)
2274              if (ch->action == KBREQUEST)
2275                    ch->flags &= ~XECUTED;
2276        DELSET(got_signals, SIGWINCH);
2277    }
2278
2279    if (ISMEMBER(got_signals, SIGALRM)) {          对收到的 SIGALRM signal，只是删除标志了事
2280        INITDBG(L_VB, "got SIGALRM");
2281        /* The timer went off: check it out */
2282        DELSET(got_signals, SIGALRM);
2283    }
2284
2285    if (ISMEMBER(got_signals, SIGCHLD)) {          收到 init process 的子进程死亡而发来的消息
2286        INITDBG(L_VB, "got SIGCHLD");
2287        /* First set flag to 0 */
2288        DELSET(got_signals, SIGCHLD);              首先删除收到该 signal 的标志
2289
2290        /* See which child this was */
2291        for(ch = family; ch; ch = ch->next)        对 family 链表进行枚举，如果发现僵尸，则清除三个标志
2292            if (ch->flags & ZOMBIE) {
2293              INITDBG(L_VB, "Child died, PID= %d", ch->pid);
2294              ch->flags &= ~(RUNNING|ZOMBIE|WAITING);
2295              if (ch->process[0] != '+')
2296                    write_utmp_wtmp("", ch->id, ch->pid, DEAD_PROCESS, NULL);
2297            }
2298
2299    }
2300
2301    if (ISMEMBER(got_signals, SIGHUP)) {     signal SIGHUP 一般用来通知进程重读其配置文件，这里就是通知 init
2302        INITDBG(L_VB, "got SIGHUP");         process 重新读取 inittab 文件中的配置
2303  #if CHANGE_WAIT
2304        /* Are we waiting for a child? */
```

```
2305        for(ch = family; ch; ch = ch->next)
2306              if (ch->flags & WAITING) break;
2307        if (ch == NULL)
2308  #endif
2309        {
2310              /* We need to go into a new runlevel */
2311              oldlevel = runlevel;
2312  #ifdef INITLVL
2313              runlevel = read_level(0);
2314  #endif
2315              if (runlevel == 'U') {
2316                    runlevel = oldlevel;
2317                    re_exec();
2318              } else {
2319                    if (oldlevel != 'S' && runlevel == 'S') console_stty();
2320                    if (runlevel == '6' || runlevel == '0' ||
2321                        runlevel == '1') console_stty();
2322                    read_inittab();
2323                    fail_cancel();
2324                    setproctitle("init [%c]", runlevel);
2325                    DELSET(got_signals, SIGHUP);
2326              }
2327        }
2328    }
2329    if (ISMEMBER(got_signals, SIGUSR1)) {          收到用户自定义 signal，这里用来重新打开 pipe 文件/dev/initctl
2330        /*
2331         *    SIGUSR1 means close and reopen /dev/initctl
2332         */
2333        INITDBG(L_VB, "got SIGUSR1");
2334        close(pipe_fd);
2335        pipe_fd = -1;
2336        DELSET(got_signals, SIGUSR1);
2337    }
2338  }
```

继续刚才未完的 init_main()的执行。

```
2400
2401    console_init();         对系统 console 的初始化
2402
2403    if (!reload) {          前面已经分析过，在 init 1 时，该值为 0，也即会进入下面的 if 分支
2404
2405        /* Close whatever files are open, and reset the console. */
2406        close(0);           这里的一系列操作，关闭 0，1，2，调用 setsid()类似于 Linux 下 Daemon 编程的老套路。init 进程也
2407        close(1);           是 daemon process
2408        close(2);
2409        console_stty();
2410        setsid();
2411
2412        /*
2413         *    Set default PATH variable.
2414         */
2415        putenv(PATH_DFL);       设置 init 1 执行时的环境变量中的 PATH="PATH=/bin:/usr/bin:/sbin:/usr/sbin"
2416
2417        /*
2418         *    Initialize /var/run/utmp (only works if /var is on
2419         *    root and mounted rw)
2420         */
2421        (void) close(open(UTMP_FILE, O_WRONLY|O_CREAT|O_TRUNC, 0644));     创建/var/run/utmp 文件
2422
2423        /*
2424         *    Say hello to the world
2425         */
2426        initlog(L_CO, bootmsg, "booting");
2427
2428        /*
2429         *    See if we have to start an emergency shell.
2430         */
2431        if (emerg_shell) {    如果在命令行上指定启动紧急 shell，但 init 1 的命令行是空的，所以这里的 if 分支不会执行到
2432              SETSIG(sa, SIGCHLD, SIG_DFL, SA_RESTART);
2433              if (spawn(&ch_emerg, &f) > 0) {
2434                    while((rc = wait(&st)) != f)
2435                          if (rc < 0 && errno == ECHILD)
```

```
2436                                break;
2437            }
2438            SETSIG(sa, SIGCHLD,  chld_handler, SA_RESTART);
2439        }
2440
2441        /*
2442         *    Start normal boot procedure.
2443         */
2444        runlevel = '#';          表示现在还不知道 init 1 将要进入什么运行级别
2445        read_inittab();          读取/etc/inittab 中的设置，非常关键的函数，见对该函数的注解
2446
2447    } else {  re-exec 模式，即 init process 运行期间重读 inittab 配置文件的处理，init 1 不满足该条件，所以跳到 L2455
2448        /*
2449         *    Restart: unblock signals and let the show go on
2450         */
2451        initlog(L_CO, bootmsg, "reloading");
2452        sigfillset(&sgt);
2453        sigprocmask(SIG_UNBLOCK, &sgt, NULL);
2454    }
2455    start_if_needed();     枚举在 family 链表中的代表/etc/inittab 每一行的 action，并启动它们。该函数见辅助函数分析。
                               作为系统启动阶段，运行 inittab 文件中指定的命令行。
2456
2457    while(1) {       进入 init process 的主循环，init process 以后就在该无限循环中打转，永不出来。这实际上是 init 3，
2458                     也即作为 daemon process 运行的 init 的工作了。
2459       /* See if we need to make the boot transitions. */
2460       boot_transitions();
2461       INITDBG(L_VB, "init_main: waiting..");
2462
2463       /* Check if there are processes to be waited on. */
2464       for(ch = family; ch; ch = ch->next)
2465        if ((ch->flags & RUNNING) && ch->action != BOOT) break;



2466
2467  #if CHANGE_WAIT
```

```
2468        /* Wait until we get hit by some signal. */
2469        while (ch != NULL && got_signals == 0) {    daemon process 检查是否有关心的事件发生。
2470         if (ISMEMBER(got_signals, SIGHUP)) {       检查发生了 SIGHUP，即让 init process 重读 inittab 文件的事件吗
2471               /* See if there are processes to be waited on. */
2472               for(ch = family; ch; ch = ch->next)        标志为“ wait ”的配置行于“ boot ”配置行类似，不能与其他
2473                     if (ch->flags & WAITING) break;      配置行并发执行
2474         }
2475         if (ch != NULL) check_init_fifo();
2476        }
2477  #else /* CHANGE_WAIT */
2478        if (ch != NULL && got_signals == 0) check_init_fifo();
2479  #endif /* CHANGE_WAIT */
```

check_init_fifo()首先建立 init 2 与 init 3 之间沟通的渠道“ /dev/initctl ”pipe。

```
2480
2481        /* Check the 'failing' flags */
2482        fail_check();
2483
2484        /* Process any signals. */
2485        process_signals();          处理被记录下来的 signal，即当 init process 收到 signal，并纪录在 got_signals 变量
2486                                    后，要在 init process 的主循环中才能执行，所以 init process 对相关 signal 的真正
                                    处理并不是实时的，即受到 signal 就处理，而是要在 L2457 的循环中查询到后才能执行
2487        /* See what we need to start up (again) */
2488        start_if_needed();          有可能在 family 链表中的 node 状态已经改变，所以重新枚举整个链表，看是否有本来不能运行
2489      }                             的动作(action)可以执行了
2490      /*NOTREACHED*/
2491  }
```

下面是 init process 及其重要的读取配置文件 inittab 并生成自己管理的数据结构。Inittab 配置文件的格式已经在上面说明过。、每一个 init process 管理的进程都用如下结构来管理：

```
/* Information about a process in the in-core inittab */
typedef struct _child_ {
  int flags;                    /* Status of this entry */
  int exstat;                   /* Exit status of process */
```

```
  int pid;                    /* Pid of this process */
  time_t tm;                  /* When respawned last */
  int count;                  /* Times respawned in the last 2 minutes */
  char id[8];                 /* Inittab id (must be unique) */
  char rlevel[12];            /* run levels */
  int action;                 /* what to do (see list below) */
  char process[128];          /* The command line */
  struct _child_ *new;        /* New entry (after inittab re-read) */
  struct _child_ *next;       /* For the linked list */
} CHILD;
```

该结构中的原有注释写得挺详细的，这里补充几点。
Inittab 中的配置行大约如下：

3:2345:respawn:/sbin/mingetty tty3
这里 init process 就会有一个 CHILD 与其对应，用上面的信息来填充该结构中的某些 field。

id[8]	rlevel[12]	action	process[128]
3	2345	Respawn	/sbin/mingetty tty3

而 flags 反映该 process 的状态
exstat 是该 process 在退出后的退出码，也就是提供给 exit()系统调用的参数
pid 是该 process 的 process identifier
tm 对 respawn 类型的 process 才有意义，即其被 init process respawn 时的时间戳
count 对 respawn 与 ondemanded 型的 process 有意义，即其在最近 2 分钟内被启动(spawn)的次数。

init process 把它管理的 process 通过 next 来链接在由全局变量 family 指向的链表中。

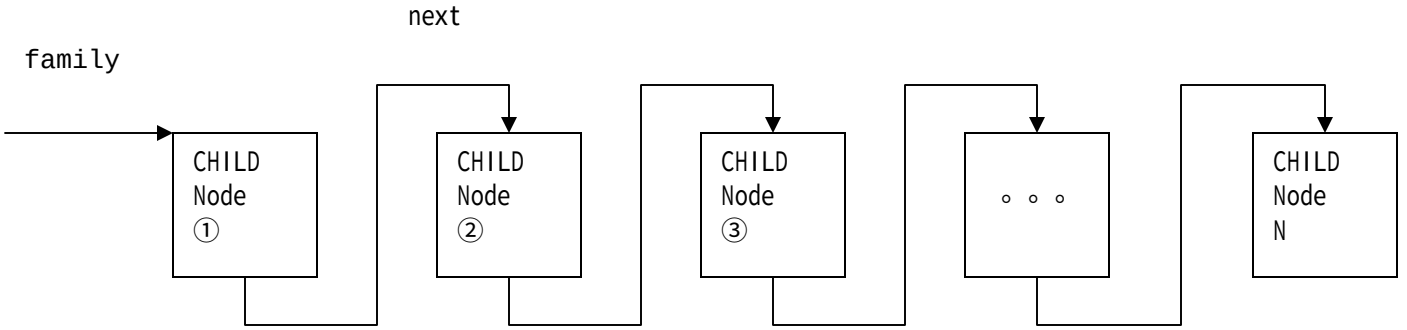

/etc/inittab
配置文件

```
id:5:initdefault:                CHILD Node ①

si::sysinit:/etc/rc.d/rc.sysinit
                                 CHILD Node ②
l0:0:wait:/etc/rc.d/rc 0         CHILD Node ③
l1:1:wait:/etc/rc.d/rc 1
l2:2:wait:/etc/rc.d/rc 2
l3:3:wait:/etc/rc.d/rc 3
l4:4:wait:/etc/rc.d/rc 4
l5:5:wait:/etc/rc.d/rc 5
l6:6:wait:/etc/rc.d/rc 6


ca::ctrlaltdel:/sbin/shutdown -t3 -r now

3:2345:respawn:/sbin/mingetty tty3  CHILD Node N
```

read_inittab () 就是把 inittab 文件中的每一行生成 family 链表中的一个节点（node）,并用该行上的信息来填充该节点

```
1108  void read_inittab(void)
1109  {
1110    FILE            *fp;              /* The INITTAB file */
1111    CHILD           *ch, *old, *i;    /* Pointers to CHILD structure */
1112    CHILD           *head = NULL;     /* Head of linked list */
1113  #ifdef INITLVL
1114    struct stat     st;               /* To stat INITLVL */
1115  #endif
1116    sigset_t  nmask, omask;           /* For blocking SIGCHLD. */
1117    char            buf[256];         /* Line buffer */
1118    char            err[64];          /* Error message. */
1119    char            *id, *rlevel,
1120              *action, *process;      /* Fields of a line */
1121    char            *p;
1122    int       lineNo = 0;       /* Line number in INITTAB file */
1123    int       actionNo;         /* Decoded action field */
1124    int       f;                /* Counter */
1125    int       round;            /* round 0 for SIGTERM, 1 for SIGKILL */
1126    int       foundOne = 0;     /* No killing no sleep */
1127    int       talk;             /* Talk to the user */
1128    int       done = 0;         /* Ready yet? */
1129
1130  #if DEBUG
1131    if (newFamily != NULL) {
1132        INITDBG(L_VB, "PANIC newFamily != NULL");
1133        exit(1);
1134    }
1135    INITDBG(L_VB, "Reading inittab");
1136  #endif
1137
1138    /*
1139     *  Open INITTAB and real line by line.
1140     */
1141    if ((fp = fopen(INITTAB, "r")) == NULL)            打开/etc/inittab 文件
1142        initlog(L_VB, "No inittab file found");
1143
```

```
1144    while(!done) {              每循环一次即处理 inittab 中一行，构造 newFamily 链表。注意是 newFamily 链表，不是
1145        /*                      family 链表
1146         *    Add single user shell entry at the end.
1147         */
1148        if (fp == NULL || fgets(buf, sizeof(buf), fp) == NULL) {
1149            done = 1;           已经处理完毕，退出循环
1150            /*
1151             *    See if we have a single user entry.
1152             */
1153            for(old = newFamily; old; old = old->next)            ？？？
1154                if (strpbrk(old->rlevel, "S")) break;
1155            if (old == NULL)
1156                snprintf(buf, sizeof(buf), "~~:S:wait:%s\n", SULOGIN);
1157            else
1158                continue;
1159        }
1160        lineNo++;
1161        /*
1162         *    Skip comments and empty lines
1163         */
1164        for(p = buf; *p == ' ' || *p == '\t'; p++)      忽略前导空白字符
1165            ;
1166        if (*p == '#' || *p == '\n') continue;          以“#”开头的为注释，忽略
1167
1168        /*
1169         *    Decode the fields
1170         */
      分解 id:runlevels:action:process 中的 4 部分
1171        id =      strsep(&p, ":");    由于文件中的配置行的各部分用“：”分割，所以这里通过 strsep 来分别提取各部分内
1172        rlevel =  strsep(&p, ":");    容
1173        action =  strsep(&p, ":");
1174        process = strsep(&p, "\n");
1175
      从下面的代码可以看到在 init manual 中没有标明的限制，比如命令行的长度不能太长，超过 127 个字符等
1176        /*
1177         *    Check if syntax is OK. Be very verbose here, to
```

```
1178          *    avoid newbie postings on comp.os.linux.setup :)
1179          */
1180         err[0] = 0;
1181         if (!id || !*id) strcpy(err, "missing id field");
1182         if (!rlevel)     strcpy(err, "missing runlevel field");
1183         if (!process)    strcpy(err, "missing process field");
1184         if (!action || !*action)
1185                     strcpy(err, "missing action field");
1186         if (id && strlen(id) > sizeof(utproto.ut_id))
1187               sprintf(err, "id field too long (max %d characters)",
1188                     (int)sizeof(utproto.ut_id));
1189         if (rlevel && strlen(rlevel) > 11)
1190               strcpy(err, "rlevel field too long (max 11 characters)");
1191         if (process && strlen(process) > 127)
1192               strcpy(err, "process field too long");
1193         if (action && strlen(action) > 32)
1194               strcpy(err, "action field too long");
1195         if (err[0] != 0) {
1196               initlog(L_VB, "%s[%d]: %s", INITTAB, lineNo, err);
1197               INITDBG(L_VB, "%s:%s:%s:%s", id, rlevel, action, process);
1198               continue;
1199         }
1200
1201         /*
1202          *    Decode the "action" field
1203          */
```

init 允许的 action 类型都记录在 actions[]数组中，这里通过比较字符串来把其转换成数字型 identifier

```
1204         actionNo = -1;
1205         for(f = 0; actions[f].name; f++)
1206               if (strcasecmp(action, actions[f].name) == 0) {
1207                     actionNo = actions[f].act;
1208                     break;
1209               }
1210         if (actionNo == -1) {
1211               initlog(L_VB, "%s[%d]: %s: unknown action field",
1212                     INITTAB, lineNo, action);
1213               continue;
```

碰到非法的 action(不在 actions[]数组中的)则忽略

```
1214        }
1215
1216        /*
1217         *    See if the id field is unique
1218         */
```

配置行中的第一部分是所谓 identifier，必须唯一，但命名好像没有任何规定，可任意。已经处理过的配置行都被记录入 CHILD 的链表节点中，这里在处理当前行时检查一下已有节点中是否有与当前行的 id 相同的，如果有，则不是忽略该行，而是停止继续处理/etc/inittab 文件，可见 id 的唯一性是至关重要的

```
1219        for(old = newFamily; old; old = old->next) {
1220              if(strcmp(old->id, id) == 0 && strcmp(id, "~~")) {
1221                    initlog(L_VB, "%s[%d]: duplicate ID field \"%s\"",
1222                          INITTAB, lineNo, id);
1223                    break;
1224              }
1225        }
1226        if (old) continue;
1227
1228        /*
1229         *    Allocate a CHILD structure
1230         */
1231        ch = imalloc(sizeof(CHILD));        为当前配置行分配一个 CHILD node
1232
1233        /*
1234         *    And fill it in.          用当前配置行中的信息来填充 CHILD node
1235         */
1236        ch->action = actionNo;         action 类型
1237        strncpy(ch->id, id, sizeof(utproto.ut_id) + 1); /* Hack for different libs. */该行的唯一标示符
1238        strncpy(ch->process, process, sizeof(ch->process) - 1);      该行是要执行的命令行
1239        if (rlevel[0]) {               填上 run level
1240              for(f = 0; f < sizeof(rlevel) - 1 && rlevel[f]; f++) {
1241                    ch->rlevel[f] = rlevel[f];
1242                    if (ch->rlevel[f] == 's') ch->rlevel[f] = 'S';
1243              }
1244              strncpy(ch->rlevel, rlevel, sizeof(ch->rlevel) - 1);
```

```
1245        } else {    如果没有写 run  level，则表示所有 run  level 都要执行该行的 process 部分
1246              strcpy(ch->rlevel, "0123456789");
1247              if (ISPOWER(ch->action))
1248                    strcpy(ch->rlevel, "S0123456789");
1249        }
      下面是对 action 的处理
1250        /*
1251         *    We have the fake runlevel '#' for SYSINIT  and
1252         *    '*' for BOOT and BOOTWAIT.
1253         */

      从上面的注释看，SYSINIT  action 用'#'表示，而 BOOT  action 用'*'表示。而真正合法的 run  level 是 0 到 9 加上'S'
      “ # ”与“ * ”表示在任何 run  level 都要执行，另外 SYSINIT 的优先级是最高的，所以它应该比 BOOT 中的 action 先执行

1254        if (ch->action == SYSINIT) strcpy(ch->rlevel, "#");
1255        if (ch->action == BOOT || ch->action == BOOTWAIT)
1256              strcpy(ch->rlevel, "*");
1257
1258        /*
1259         *    Now add it to the linked list. Special for powerfail.
1260         */

在从/etc/inittab 中读取配置行并生成的链表的头为 newFamily。如果是系统的启动阶段，family 所表示的链表自然为空，如果只是通过运
行 init 来切换 run  level 等，则 family 记录的链表非空，也就是当前 init 通过上次读取/etc/inittab 后生成的链表。

1261        if (ISPOWER(ch->action)) {    如果 action 属于这几种（POWERWAIT，POWERFAIL，POWEROKWAIT，
1262                                      POWERFAILNOW，CTRLALTDEL），即与电源相关与用户按了 Ctrl+Alt+Del 键

1263              /*
1264               *    Disable by default
1265               */
1266              ch->flags |= XECUTED;   设置不执行标志。在 startup()种如果检测到该配置行所代表的 action 的 flag 设置
1267                                      了 XECUTED，则忽略对该行的处理。这可以理解，因为符合 ISPOWER()的 action 都不是
1268              /*                      在正常情况下需要运行的，只有对应的事件确实发生了，才需要执行。比如，如果用户从
                                        来没有按下过 Ctrl+Alt+Del 键，自然根本不需要执行/etc/inittab 中 CTRLALTDEL
```

Action 指定的 process。所以在默认情况下它是 disable 的(通过设置 XECUTED 标志)。当检测到按下 Ctrl+Alt+Del 键后，才需要 enable。
并且上述 action 被插在 family 链表的前面，这样如果它们需要执行的话，将先得到执行。可以理解，因为这几个动作都比较严重，所以优先权较高

```
1269                *    Tricky: insert at the front of the list..
1270                */
1271               old = NULL;
1272               for(i = newFamily; i; i = i->next) {
1273                     if (!ISPOWER(i->action)) break;
1274                     old = i;
1275               }
1276               /*
1277                *    Now add after entry "old"
1278                */
1279               if (old) {
1280                     ch->next = i;
1281                     old->next = ch;
1282                     if (i == NULL) head = ch;
1283               } else {
1284                     ch->next = newFamily;
1285                     newFamily = ch;
1286                     if (ch->next == NULL) head = ch;
1287               }
1288         } else {  其他的 action 都插到尾部，KBREQUEST 默认是不执行的。从 init manual 上看 SIGWINCH  signal 会触发该动作
1289               /*
1290                *    Just add at end of the list
1291                */
1292               if (ch->action == KBREQUEST) ch->flags |= XECUTED;
1293               ch->next = NULL;
1294               if (head)
1295                     head->next = ch;
1296               else
1297                     newFamily = ch;
1298               head = ch;
1299         }
1300
```

```
1301        /*
1302         *    Walk through the old list comparing id fields
1303         */
1304        for(old = family; old; old = old->next)
1305              if (strcmp(old->id, ch->id) == 0) {
1306                    old->new = ch;
1307                    break;
1308              }
            到这里处理一行结束
1309    }
1310    /*
1311     *  We're done.
1312     */
1313    if (fp) fclose(fp);    关闭/etc/inittab 文件，对于 init 1，实际上基本上到此为止。
1314                           下面都是作为 daemon 进程的 init 即 init 3 的处理。我觉得整个代码应该整理得更清晰一点，由
                               kernel 启动 init(init 1)与作为 daemon 进程运行的 init（init 3）的逻辑应该分开，而不要像现
                               在一样，绞和在一块，比较乱

1315    /*
1316     *  Loop through the list of children, and see if they need to
1317     *  be killed.
1318     */
1319
1320    INITDBG(L_VB, "Checking for children to kill");
1321    for(round = 0; round < 2; round++) {
1322      talk = 1;
1323      for(ch = family; ch; ch = ch->next) { 由于是系统启动阶段运行 init（init 1），则这时的 family 链表为空，应该
1324        ch->flags &= ~KILLME;                不进入循环，则这时 round = 0，talk = 1，foundOne = 0，代码应该跳转
1325                                             到 L1393 行执行。



1326        /*
1327         *    Is this line deleted?
1328         */
1329        if (ch->new == NULL) ch->flags |= KILLME;
```

```
1330
1331        /*
1332         *    If the entry has changed, kill it anyway. Note that
1333         *    we do not check ch->process, only the "action" field.
1334         *    This way, you can turn an entry "off" immediately, but
1335         *    changes in the command line will only become effective
1336         *    after the running version has exited.
1337         */
1338        if (ch->new && ch->action != ch->new->action) ch->flags |= KILLME;
1339
1340        /*
1341         *    Only BOOT processes may live in all levels
1342         */
1343        if (ch->action != BOOT &&
1344            strchr(ch->rlevel, runlevel) == NULL) {
1345              /*
1346               *    Ondemand procedures live always,
1347               *    except in single user
1348               */
1349              if (runlevel == 'S' || !(ch->flags & DEMAND))
1350                    ch->flags |= KILLME;
1351        }
1352
1353        /*
1354         *    Now, if this process may live note so in the new list
1355         */
1356        if ((ch->flags & KILLME) == 0) {
1357              ch->new->flags  = ch->flags;
1358              ch->new->pid    = ch->pid;
1359              ch->new->exstat = ch->exstat;
1360              continue;
1361        }
1362
1363
1364        /*
1365         *    Is this process still around?
1366         */
```

```
1367        if ((ch->flags & RUNNING) == 0) {
1368              ch->flags &= ~KILLME;
1369              continue;
1370        }
1371        INITDBG(L_VB, "Killing \"%s\"", ch->process);
1372        switch(round) {
1373              case 0: /* Send TERM signal */
1374                    if (talk)
1375                          initlog(L_CO,
1376                                "Sending processes the TERM signal");
1377                    kill(-(ch->pid), SIGTERM);
1378                    foundOne = 1;
1379                    break;
1380              case 1: /* Send KILL signal and collect status */
1381                    if (talk)
1382                          initlog(L_CO,
1383                                "Sending processes the KILL signal");
1384                    kill(-(ch->pid), SIGKILL);
1385                    break;
1386        }
1387        talk = 0;
1388
1389      }
1390      /*
1391       *      See if we have to wait 5 seconds
1392       */
1393      if (foundOne && round == 0) {   对于系统启动阶段运行的 init（init 1），round = 0,但 foundOne = 0，所以
1394        /*                            将不会进入该 if 分支，跳转到 L1419 执行
1395         *    Yup, but check every second if we still have children.
1396         */
1397        for(f = 0; f < sltime; f++) {
1398              for(ch = family; ch; ch = ch->next) {
1399                    if (!(ch->flags & KILLME)) continue;
1400                    if ((ch->flags & RUNNING) && !(ch->flags & ZOMBIE))
1401                          break;
1402              }
1403              if (ch == NULL) {
```

```
1404                    /*
1405                     *    No running children, skip SIGKILL
1406                     */
1407                    round = 1;
1408                    foundOne = 0; /* Skip the sleep below. */
1409                    break;
1410            }
1411            do_sleep(1);
1412        }
1413      }
1414    }
1415
1416    /*
1417     *  Now give all processes the chance to die and collect exit statuses.
1418     */
1419    if (foundOne) do_sleep(1);              init 1 运行时，foundOne = 0，所以不会睡眠一秒
1420    for(ch = family; ch; ch = ch->next)     这时的 family 链表为空，不进入该循环，跳转到 L1437 执行
1421        if (ch->flags & KILLME) {
1422            if (!(ch->flags & ZOMBIE))
1423                initlog(L_CO, "Pid %d [id %s] seems to hang", ch->pid,
1424                        ch->id);
1425            else {
1426                INITDBG(L_VB, "Updating utmp for pid %d [id %s]",
1427                        ch->pid, ch->id);
1428                ch->flags &= ~RUNNING;
1429                if (ch->process[0] != '+')
1430                  write_utmp_wtmp("", ch->id, ch->pid, DEAD_PROCESS, NULL);
1431            }
1432        }
1433
1434    /*
1435     *  Both rounds done; clean up the list.
1436     */
1437    sigemptyset(&nmask);
1438    sigaddset(&nmask, SIGCHLD);
1439    sigprocmask(SIG_BLOCK, &nmask, &omask);
1440    for(ch = family; ch; ch = old) {            init 1 运行时，family 链表还为空，所以不进入循环，跳转到 L1444
```

```
1441        old = ch->next;
1442        free(ch);
1443    }
1444    family = newFamily;                newFamily 中就是在本函数头上读入的当前的/etc/inittab 配置行生成的链表，现在才
                                           把它赋给 family

1445    for(ch = family; ch; ch = ch->next) ch->new = NULL;
1446    newFamily = NULL;
1447    sigprocmask(SIG_SETMASK, &omask, NULL);
1448
1449  #ifdef INITLVL
1450    /*
1451     *  Dispose of INITLVL file.
1452     */ 删除/etc/initrunlvl，根据是文件还是 symbol link，删除方式是不一样的
1453    if (lstat(INITLVL, &st) >= 0 && S_ISLNK(st.st_mode)) {    检查/etc/initrunlvl 是文件还是 symbol link
1454        /*
1455         *    INITLVL is a symbolic link, so just truncate the file.
1456         */
1457        close(open(INITLVL, O_WRONLY|O_TRUNC));
1458    } else {
1459        /*
1460         *    Delete INITLVL file.
1461         */
1462        unlink(INITLVL);
1463    }
1464  #endif
1465  #ifdef INITLVL2
1466    /*
1467     *  Dispose of INITLVL2 file.
1468     */ 删除/var/log/initrunlvl，根据是文件还是 symbol link，删除方式是不一样的
1469    if (lstat(INITLVL2, &st) >= 0 && S_ISLNK(st.st_mode)) {
1470        /*
1471         *    INITLVL2 is a symbolic link, so just truncate the file.
1472         */
1473        close(open(INITLVL2, O_WRONLY|O_TRUNC));
1474    } else {
1475        /*
```

```
1476          *    Delete INITLVL2 file.
1477          */
1478         unlink(INITLVL2);
1479       }
1480  #endif
1481  }
```

从上面的代码分析可看出，对 init 1 而言，read_inittab()的处理到 L1313 行实际上就结束了，剩下的处理都是为作为 daemon 运行的 init 3 来服务的（将在分析 init 3 的运行时对其详细分析），它的功能就是读取 inittab 配置文件中的行，然后构建 CHILD node 构成的 family 链表，即把硬盘上的 init process 要实现的功能的配置转换成内存中的实现配置，以便 start_if_needed()根据该内存中的配置来执行。

```
1483  /*
1484   *    Walk through the family list and start up children.
1485   *    The entries that do not belong here at all are removed
1486   *    from the list.
1487   */
```

从上面的注释可看出本函数的功能就是枚举 family 链表，然后执行每个节点上指定的动作，即执行 CHILD 节点中 process[128]中指定的命令行

```
1488  void start_if_needed(void)
1489  {
1490        CHILD *ch;        /* Pointer to child */
1491        int delete;       /* Delete this entry from list? */
1492
1493        INITDBG(L_VB, "Checking for children to start");
```

通过 read_inittab()，所有/etc/inittab 配置文件中的合法的每一行都被用 family 链表中的一个 node 表示，现在枚举该链表

```
1494
1495        for(ch = family; ch; ch = ch->next) {
1496
1497  #if DEBUG
1498              if (ch->rlevel[0] == 'C') {
1499                    INITDBG(L_VB, "%s: flags %d", ch->process, ch->flags);
```

```
1500                }
1501  #endif
1502
1503                /* Are we waiting for this process? Then quit here. */
1504                if (ch->flags & WAITING) break;      WAITING 标志表示 init process 必须等待该启动的子进程退出后才能
1505                                                     继续下面的工作，即可能后面的工作依赖于当前的 process，所以必须
                                                         等待其完成后才能开始新的工作

1506                /* Already running? OK, don't touch it */
1507                if (ch->flags & RUNNING) continue;   该 node 所代表的 process 处于 running 状态，当然不需要做什么
1508
1509                /* See if we have to start it up */
1510                delete = 1;                          默认置删除标志
1511                if (strchr(ch->rlevel, runlevel) ||        runlevel 变量记录着当前的 run level,对于在该 node 中包含
                                                           了当前运行级别的 node 或者是 DEMAND(按需启动)或者
                                                           “ #*Ss ”（代表 sysinit，boot 等）则启动该节点所代表的
                                                           action。(“ #*Ss ”部分 run level 都要执行)

1512                    ((ch->flags & DEMAND) && !strchr("#*Ss", runlevel))) {
1513                      startup(ch);       与执行该命令，具体间辅助函数介绍
1514                      delete = 0;        启动了，当然不需要删除该 node，则 clear 该标志
1515                }
1516
1517                if (delete) {      没有启动，表明该 node 不符合合法的启动 level
1518                      /* FIXME: is this OK? */
1519                      ch->flags &= ~(RUNNING|WAITING);
1520                      if (!ISPOWER(ch->action) && ch->action != KBREQUEST)
1521                            ch->flags &= ~XECUTED;
1522                      ch->pid = 0;
1523                } else
1524                      /* Do we have to wait for this process? */
1525                      if (ch->flags & WAITING) break;
1526          }
1527          /* Done. */
1528    }
```

辅助函数介绍

startup函数

```
1063  /*
1064   *    Start a child running!
1065   */
```

执行 CHILD 节点所代表的配置行上的命令行，比如 l2:2:wait:/etc/rc.d/rc 2 配置行，就是执行这里的“ /etc/rc.d/rc 2 ”，“ /etc/rc.d/rc ”是执行的可执行脚本，而“ 2 ”是该脚本的参数

```
1066  void startup(CHILD *ch)
1067  {
1068        /*
1069         *    See if it's disabled
1070         */
1071        if (ch->flags & FAILING) return;
1072
1073        switch(ch->action) {
1074
1075              case SYSINIT:
1076              case BOOTWAIT:
1077              case WAIT:
1078              case POWERWAIT:
1079              case POWERFAILNOW:
1080              case POWEROKWAIT:
1081              case CTRLALTDEL:
1082                    if (!(ch->flags & XECUTED)) ch->flags |= WAITING;
```

对于上面的几种 action，如果已经允许执行（没有设 XECUTED 标志），则会设置 WAITING 标志，由于这里没有 break,所以将执行 L1091 行的 spawn()，即启动该命令。之所以这里要设 WAITING 标志，是因为上面的类型的动作都具有排它性。

```
1083              case KBREQUEST:
```

```
1084                case BOOT:
1085                case POWERFAIL:
1086                case ONCE:
1087                      if (ch->flags & XECUTED) break;      XECUTED 是暂时禁止执行的标志
1088                case ONDEMAND:
1089                case RESPAWN:
1090                      ch->flags |= RUNNING;
1091                      if (spawn(ch, &(ch->pid)) < 0) break;      启动该脚本，ch->pid 为返回的 pid，spawn()见辅助
1092                      /*                                         函数说明
1093                       *    Do NOT log if process field starts with '+'
1094                       *    FIXME: that's for compatibility with *very*
1095                       *    old getties - probably it can be taken out.
1096                       */
1097                      if (ch->process[0] != '+')
1098                            write_utmp_wtmp("", ch->id, ch->pid,
1099                                  INIT_PROCESS, "");
1100                      break;
1101          }
1102    }
```

spawn函数

该函数是真正的去 launch 在 CHILD 中指定的命令行（process[128]）

```
823   /*
824    *    Fork and execute.
825    *
826    *    This function is too long and indents too deep.
827    *
828    */   启动 ch 所对应的配置行上的命令，*res 为启动的这个新进程的 pid
```

下面以启动/etc/inittab 中的如下配置行为例：

```
si::sysinit:/etc/rc.d/rc.sysinit
```

则 CHILD *ch 的内容如下

```
ch->id = “ si ”
```

```
ch->rlevel = 0123456S
ch->action = SYSINIT
ch->process[128] = /etc/rc.d/rc.sysinit

829   int spawn(CHILD *ch, int *res)
830   {
831     char *args[16];       /* Argv array */
832     char buf[136];        /* Line buffer */
833     int f, st, rc;        /* Scratch variables */
834     char *ptr;            /* Ditto */
835     time_t t;             /* System time */
836     int oldAlarm;            /* Previous alarm value */
837     char *proc = ch->process;   /* Command line */
838     pid_t pid, pgrp;         /* child, console process group. */
839     sigset_t nmask, omask;      /* For blocking SIGCHLD */
840     struct sigaction sa;
841
842     *res = -1;
843     buf[sizeof(buf) - 1] = 0;
844
845     /* Skip '+' if it's there */
846     if (proc[0] == '+') proc++;
847
848     ch->flags |= XECUTED;       当前正要启动的 ch，设置禁止启动标志，防止再进入
849
```

下面是根据不同的情况来拼装命令行字符串

init process 将监控 action 为“ RESPAWN ”与“ ONDEMAND ”类型的命令，如果其不运行了，则要启动它，使它在整个系统运行期间一直运行。下面的处理是为了防止由于该命令的不正常死亡，导致 init process 在短时间内太过频繁的启动。如果在 2 分钟内启动超过 10 次，则先要把该命令凉在一边 5 分钟，然后再启动。

```
850     if (ch->action == RESPAWN || ch->action == ONDEMAND) {
851         /* Is the date stamp from less than 2 minutes ago? */
852         time(&t);    取得当前系统时间，ch->tm 为上次该 process 启动(spawn)时的时间
853         if (ch->tm + TESTTIME > t) {  TESTTIME 为 2 分钟，即从上次启动到现在没到 2 分钟，则累加被启动的次数
```

```
854                 ch->count++;
855         } else {
856                 ch->count = 0;          超过 2 分钟则重新计数在 2 分钟内被启动次数
857                 ch->tm = t;             重设时间
858         }
859
860         /* Do we try to respawn too fast? */
861         if (ch->count >= MAXSPAWN) {  启动得太频繁了，2 分钟内被启动了 10 次（说明该程序在 2 分钟内死了超过 10 次）设置
862                                       FAILING 标志
863           initlog(L_VB,
864               "Id \"%s\" respawning too fast: disabled for %d minutes",
865               ch->id, SLEEPTIME / 60);
866           ch->flags &= ~RUNNING;
867           ch->flags |= FAILING;       临时设上 FAILING 标志
868
869           /* Remember the time we stopped */
870           ch->tm = t;
871
872           /* Try again in 5 minutes */
873           oldAlarm = alarm(0);        取消 alarm，返回剩余时间
874           if (oldAlarm > SLEEPTIME || oldAlarm <= 0) oldAlarm = SLEEPTIME;      最多暂时 disable 5 分钟
875           alarm(oldAlarm);            启动 alarm，最多等待 5 分钟
876           return(-1);           没有启动就返回了（-1）
877         }
878     }
879
880     /* See if there is an "initscript" (except in single user mode). */
881     if (access(INITSCRIPT, R_OK) == 0 && runlevel != 'S') {   文件/etc/initscript 可读吗，实际上就是是否存在
882         /* Build command line using "initscript" */          如果/etc/initscript 存在，并且当前 run level
883         args[1] = SHELL;                                     不为 Single Mode(单用户模式)
884         args[2] = INITSCRIPT;
885         args[3] = ch->id;                   为 L1045 行的 execvp()调用准备参数
886         args[4] = ch->rlevel;               这里执行的命令大致是这样的
887         args[5] = "unknown";                execvp(“/bin/sh”,args + 1),而 args[1]=“/bin/sh”
888         for(f = 0; actions[f].name; f++) {                              args[2] = “/etc/initscript”
```

```
889                 if (ch->action == actions[f].act) {                          args[3] = “ si ”
890                       args[5] = actions[f].name;                             args[4] = “ 0123456S ”
891                       break;                                          args[5]=“  sysinit ”
892                 }                                                     args[6]=“ /etc/rc.d/rc.sysinit ”
893           }                                                           args[7]= NULL
894           args[6] = proc;
895           args[7] = NULL;
896       } else if (strpbrk(proc, "~`!$^&*()=|\\{}[];\"'<>?")) {   在命令行上查找以“ ~`!$^&*()=|\\{}[];\"'<>? ”
897       /* See if we need to fire off a shell for this command */ 中任何一个字符匹配的第一个字符的位置
898           /* Give command line to shell */
899           args[1] = SHELL;              运行/bin/sh –c exec /etc/rc.d/rc.sysinit
900           args[2] = "-c";
901           strcpy(buf, "exec ");
902           strncat(buf, proc, sizeof(buf) - strlen(buf) - 1);
903           args[3] = buf;
904           args[4] = NULL;
905       } else {
906           /* Split up command line arguments */
907           buf[0] = 0;
908           strncat(buf, proc, sizeof(buf) - 1);
909           ptr = buf;
910           for(f = 1; f < 15; f++) {
911                 /* Skip white space */
912                 while(*ptr == ' ' || *ptr == '\t') ptr++;
913                 args[f] = ptr;
914
915                 /* May be trailing space.. */
916                 if (*ptr == 0) break;
917
918                 /* Skip this `word' */
919                 while(*ptr && *ptr != ' ' && *ptr != '\t' && *ptr != '#')
920                       ptr++;
921
922                 /* If end-of-line, break */
923                 if (*ptr == '#' || *ptr == 0) {
924                       f++;
```

```
925                     *ptr = 0;
926                     break;
927                }
928                /* End word with \0 and continue */
929                *ptr++ = 0;
930           }
931           args[f] = NULL;
932      }
933      args[0] = args[1];
934      while(1) {
935           /*
936            *    Block sigchild while forking.
937            */
938           sigemptyset(&nmask);
939           sigaddset(&nmask, SIGCHLD);
940           sigprocmask(SIG_BLOCK, &nmask, &omask);

              按 daemon 进程方式来启动进程

941
942           if ((pid = fork()) == 0) {    init process 用它的子进程来运行启动的命令
943
944                close(0);         关闭 file handle 0,1,2, 即 STDIN, STDOUT, STDERR
945                close(1);
946                close(2);
947                if (pipe_fd >= 0) close(pipe_fd);
948
949                sigprocmask(SIG_SETMASK, &omask, NULL);
950
951                /*
952                 *    In sysinit, boot, bootwait or single user mode:
953                 *    for any wait-type subprocess we _force_ the console
954                 *    to be its controlling tty.
955                 */
956                if (strchr("*#sS", runlevel) && ch->flags & WAITING) {
957                     /*
958                      *    We fork once extra. This is so that we can
```

```
959                      *    wait and change the process group and session
960                      *    of the console after exit of the leader.
961                      */
962                     setsid();
963                     if ((f = console_open(O_RDWR|O_NOCTTY)) >= 0) {
964                           /* Take over controlling tty by force */
965                           (void)ioctl(f, TIOCSCTTY, 1);
966                           dup(f);
967                           dup(f);
968                     }
969                     if ((pid = fork()) < 0) {     fork 失败，则退出，原始 init process 当然不会退出，这里退出的
970                           initlog(L_VB, "cannot fork");       是其子进程，该进程由 L942 的 fork()产生
971                           exit(1);
972                     }
973                     if (pid > 0) {    父进程等待子进程（通过 waitpid 系统调用），也就是执行命令行的子进程返回
974                           /*
975                            *    Ignore keyboard signals etc.
976                            *    Then wait for child to exit.
977                            */
978                           SETSIG(sa, SIGINT, SIG_IGN, SA_RESTART);
979                           SETSIG(sa, SIGTSTP, SIG_IGN, SA_RESTART);
980                           SETSIG(sa, SIGQUIT, SIG_IGN, SA_RESTART);
981                           SETSIG(sa, SIGCHLD, SIG_DFL, SA_RESTART);
982
983                           while ((rc = waitpid(pid, &st, 0)) != pid)
984                                 if (rc < 0 && errno == ECHILD)
985                                       break;
986
987                           /*
988                            *    Small optimization. See if stealing
989                            *    controlling tty back is needed.
990                            */
991                           pgrp = tcgetpgrp(f);
992                           if (pgrp != getpid())
993                                 exit(0);
994
995                           /*
```

```
996                              *    Steal controlling tty away. We do
997                              *    this with a temporary process.
998                              */
999                             if ((pid = fork()) < 0) {
1000                                  initlog(L_VB, "cannot fork");
1001                                  exit(1);
1002                            }
1003                            if (pid == 0) {
1004                                  setsid();
1005                                  (void)ioctl(f, TIOCSCTTY, 1);
1006                                  exit(0);
1007                            }
1008                            while((rc = waitpid(pid, &st, 0)) != pid)
1009                                  if (rc < 0 && errno == ECHILD)
1010                                        break;
1011                            exit(0);
1012                      }
1013
1014                      /* Set ioctl settings to default ones */
1015                      console_stty();
1016
1017               } else {    子进程，命令行将在该 process 的 context 内执行
1018                      setsid();
1019                      if ((f = console_open(O_RDWR|O_NOCTTY)) < 0) {
1020                            initlog(L_VB, "open(%s): %s", console_dev,
1021                                  strerror(errno));
1022                            f = open("/dev/null", O_RDWR);
1023                      }
1024                      dup(f);     由于在 L944，L945，L946 上关闭了 0,1,2 的 handle，所以自然子进程也继承了父进程的
1025                      dup(f);     handle，即这时候没有 0，1，2 handle。这里打开 system console 作为新的 1,2,3
                                    Handle。这样当调用 console_open()打开/dev/console 时返回的 f 是 0，而 L1024 则置
                                    当前 process 的 1 handle 为/dev/console，L1025 则置当前 process 的 2 handle 为
                                    /dev/console。这样 STDIN(标准输入)，STDOUT（标准输出），STDERR（标准错误输出）都
                                    被正确设置。

                                    注意点：
```

在 Unix 环境下，启动任何程序后，其默认的前 3 个 handle(0,1,2)都已经设定，它的根源就在这里。

```
1026              }
1027
1028              /* Reset all the signals, set up environment */
1029              for(f = 1; f < NSIG; f++) SETSIG(sa, f, SIG_DFL, SA_RESTART);
                                                          把该 process 的 signal 都设为忽略

1030              environ = init_buildenv(1);          建立命令行执行的环境
1031
1032              /*
1033               *    Execute prog. In case of ENOEXEC try again
1034               *    as a shell script.
1035               */
1036              execvp(args[1], args + 1);           执行命令的调用
1037              if (errno == ENOEXEC) {
1038                    args[1] = SHELL;
1039                    args[2] = "-c";
1040                    strcpy(buf, "exec ");
1041                    strncat(buf, proc, sizeof(buf) - strlen(buf) - 1);
1042                    args[3] = buf;
1043                    args[4] = NULL;
1044                    execvp(args[1], args + 1);
1045              }
1046              initlog(L_VB, "cannot execute \"%s\"", args[1]);
1047              exit(1);
1048        }
1049        *res = pid;       把运行的命令的进程的 pid 返回给 init process
1050        sigprocmask(SIG_SETMASK, &omask, NULL);
1051
1052        INITDBG(L_VB, "Started id %s (pid %d)", ch->id, pid);
1053
1054        if (pid == -1) {
1055              initlog(L_VB, "cannot fork, retry..");
1056              do_sleep(5);
```

```
1057                continue;
1058        }
1059        return(pid);
1060    }
1061  }
```

check_init_fifo函数

```
1991  /*
1992   *    Read from the init FIFO. Processes like telnetd and rlogind can
1993   *    ask us to create login processes on their behalf.
1994   *
1995   *    FIXME:      this needs to be finished. NOT that it is buggy, but we need
1996   *          to add the telnetd/rlogind stuff so people can start using it.
1997   *          Maybe move to using an AF_UNIX socket so we can use
1998   *          the 2.2 kernel credential stuff to see who we're talking to.
1999   *
2000   */
2001  void check_init_fifo(void)
2002  {
2003    struct init_request   request;
2004    struct timeval  tv;
2005    struct stat           st, st2;
2006    fd_set          fds;
2007    int             n;
2008    int             quit = 0;
2009
2010    /*
2011     *  First, try to create /dev/initctl if not present.
2012     */
2013    if (stat(INIT_FIFO, &st2) < 0 && errno == ENOENT)          如果没有建立命名管道(name pipe)/dev/initctl,
2014        (void)mkfifo(INIT_FIFO, 0600);                         则建立之，只允许 root 用户读写
2015
2016    /*
2017     *  If /dev/initctl is open, stat the file to see if it
2018     *  is still the _same_ inode.
```

```
2019     */
2020    if (pipe_fd >= 0) {          pipe_fd 是管道/dev/initctl 的 file handle，如果 >= 0 表示已被 open
2021        fstat(pipe_fd, &st);
2022        if (stat(INIT_FIFO, &st2) < 0 ||
2023            st.st_dev != st2.st_dev ||       比较此管道是否是原始打开时的管道（通过 ino 来比较）
2024            st.st_ino != st2.st_ino) {
2025              close(pipe_fd);                如果不一致，则关闭之，这样下面的代码将再次 open 该 pipe
2026              pipe_fd = -1;                  此赋值，会造成满足 L2033 的 if 分支
2027        }
2028    }
2029
2030    /*
2031     *  Now finally try to open /dev/initctl
2032     */
2033    if (pipe_fd < 0) {                                                  pipe_fd 为/dev/initctl 的 file handle，-1
2034        if ((pipe_fd = open(INIT_FIFO, O_RDWR|O_NONBLOCK)) >= 0) {      表示还没有打开，这里非阻塞打开
2035              fstat(pipe_fd, &st);
2036              if (!S_ISFIFO(st.st_mode)) {
2037                    initlog(L_VB, "%s is not a fifo", INIT_FIFO);
2038                    close(pipe_fd);
2039                    pipe_fd = -1;
2040              }
2041        }
2042        if (pipe_fd >= 0) {
2043              /*
2044               *    Don't use fd's 0, 1 or 2.
2045               */
2046              (void) dup2(pipe_fd, PIPE_FD);       把/dev/initctl 的 file handle 设为 PIPE_FD（10）
2047              close(pipe_fd);
2048              pipe_fd = PIPE_FD;                   使得/dev/initctl 管道的 file handle 为 10(PIPE_FD)
2049
2050              /*
2051               *    Return to caller - we'll be back later.
2052               */
2053        }
2054    }
```

```
2055
```

到了这里，/dev/initctl 管道应该被正确打开

```
2056    /* Wait for data to appear, _if_ the pipe was opened. */
2057    if (pipe_fd >= 0) while(!quit) {
2058
2059        /* Do select, return on EINTR. */
2060        FD_ZERO(&fds);
2061        FD_SET(pipe_fd, &fds);
2062        tv.tv_sec = 5;                设 select 调用在该管道上最多阻塞 5 秒(timeout 为 5)
2063        tv.tv_usec = 0;
2064        n = select(pipe_fd + 1, &fds, NULL, NULL, &tv);       通过 select 调用来等待/dev/initctl 的输入，如有
2065        if (n <= 0) {                                         输入则马上返回，否则就阻塞次调用
2066              if (n == 0 || errno == EINTR) return;           即 init 3 等待 init 2 发来的 request
2067              continue;
2068        }
2069
```

从 select 调用返回，表示有进程往/dev/initctl 管道写东西，实际上就是用户通过运行 init X 把要切换到 X run level 的 request 写入该管道，造成 select 调用返回

```
2070        /* Read the data, return on EINTR. */
2071        n = read(pipe_fd, &request, sizeof(request));
```

读取写入/dev/initctl 中的 request，该 request 的格式如下：

```
struct init_request {
      int   magic;            /* Magic number                 */
      int   cmd;              /* What kind of request         */
      int   runlevel;         /* Runlevel to change to        */
      int   sleeptime;        /* Time between TERM and KILL   */
      union {
            struct init_request_bsd bsd;
            char              data[368];
      } i;
```

```
};


2072        if (n == 0) {
2073              /*
2074               *    End of file. This can't happen under Linux (because
2075               *    the pipe is opened O_RDWR - see select() in the
2076               *    kernel) but you never know...
2077               */
2078              close(pipe_fd);
2079              pipe_fd = -1;
2080              return;
2081        }
2082        if (n <= 0) {
2083              if (errno == EINTR) return;          被 signal 打断了 select （）系统调用的等待，所以继续
2084              initlog(L_VB, "error reading initrequest");
2085              continue;
2086        }
2087

      正常情况下，运行到这儿，表示确有 request 被写入/dev/initctl

2088        /*
2089         *    This is a convenient point to also try to
2090         *    find the console device or check if it changed.
2091         */
2092        console_init();
2093
2094        /*
2095         *    Process request.
2096         */
2097        if (request.magic != INIT_MAGIC || n != sizeof(request)) {   检查被写入的 request 的格式的合法性
2098              initlog(L_VB, "got bogus initrequest");
2099              continue;
2100        }
2101        switch(request.cmd) {
```

输入的 request 的格式合法，则判断要采取什么动作。即 init 2 可以向 daemon 进程的 init 3 发送的命令列表。具体 init 2 向 Init 2 发 request 的代码在 telinit()中，代码分析见对 init 2 的代码分析。

下面的全局变量 sltime 用来设定 SIGTERM 与 SIGKILL 间的间隔。当 init process 要杀死某个其管理的 process 时，先发送 SIGTERM，然后等待 sltime 秒（默认为 5 秒），再发送 SIGKILL。

```
2102              case INIT_CMD_RUNLVL:               切换到新的 run level
2103                    sltime = request.sleeptime;              取得等待时间
2104                    fifo_new_level(request.runlevel);        设置新的 run level，通过重新读取 inittab 文件来
                                                                 启动与新 run level 匹配的命令脚本
2105                    quit = 1;
2106                    break;
2107              case INIT_CMD_POWERFAIL:
2108                    sltime = request.sleeptime;
2109                    do_power_fail('F');
2110                    quit = 1;
2111                    break;
2112              case INIT_CMD_POWERFAILNOW:
2113                    sltime = request.sleeptime;
2114                    do_power_fail('L');
2115                    quit = 1;
2116                    break;
2117              case INIT_CMD_POWEROK:
2118                    sltime = request.sleeptime;
2119                    do_power_fail('O');
2120                    quit = 1;
2121                    break;
```

上面 3 个 request 都与 UPS 发来的与电源 event 有关。其处理代码都是 do_power_fail()，只不过通过不同的参数来区分状态

```
2122              case INIT_CMD_SETENV:
2123                    initcmd_setenv(request.i.data, sizeof(request.i.data));
2124                    break;
2125              case INIT_CMD_CHANGECONS:
2126                    if (user_console) {
2127                          free(user_console);
```

```
2128                        user_console = NULL;
2129                  }
2130                  if (!request.i.bsd.reserved[0])
2131                        user_console = NULL;
2132                  else
2133                        user_console = strdup(request.i.bsd.reserved);
2134                  console_init();
2135                  quit = 1;
2136                  break;
2137            default:
2138                  initlog(L_VB, "got unimplemented initrequest.");
2139                  break;
2140        }
2141    }
2142
2143    /*
2144     *  We come here if the pipe couldn't be opened.
2145     */
2146    if (pipe_fd < 0) pause();
2147
2148  }
```

init 2 的运行

init 2 是指具有 root 权限的用户通过运行 init 来切换运行级别，或设置一些 init 3(daemon 进程)在运行当中的参数，比如指定 init 在向 process 发 SIGTERM 与 SIGKILL 之间的间隔秒数(sltime)及 init 启动程序时传递的环境。

还是从 main()开始分析。

```
2597  int main(int argc, char **argv)
2598  {
2599        char              *p;
2600        int               f;
2601        int               isinit;
2602
2603        /* Get my own name */
```

```
2604        if ((p = strrchr(argv[0], '/')) != NULL)
2605              p++;
2606        else
2607              p = argv[0];
2608        umask(022);                     argv[0] = /sbin/init，则 p 指向 init
2609
2610        /* Quick check */
2611        if (geteuid() != 0) {           检查是否拥有 root 权限，运行 init 必须拥有该权限
2612              fprintf(stderr, "%s: must be superuser.\n", p);
2613              exit(1);
2614        }
2615
2616        /*
2617         *    Is this telinit or init ?
2618         */
2619        isinit = (getpid() == 1);       这时启动的 init process(init 2)的 pid 当然不可能是 1（1 正被 init 2 给占着
                                        呢），所以 isinit = 0
2620        for (f = 1; f < argc; f++) {    init 2 可以模拟 init 1 的执行，就是通过传递命令行参数-i 或—init，在手册上
                                        没有提到这一点
2621              if (!strcmp(argv[f], "-i") || !strcmp(argv[f], "--init"))
2622                    isinit = 1;
2623                    break;
2624        }
2625        if (!isinit) exit(telinit(p, argc, argv));      这时 isinit = 0，所以 init 2 的执行进入 telinit()的执
2626                                                        行
2627        /*
2628         *    Check for re-exec
2629         */

。。。。。。
```

telinit()也是 init 2 运行的主要函数。这里的参数 programe 为 init process 的可执行文件名“ init ”。

```
2502  int telinit(char *progname, int argc, char **argv)
```

```
2503  {
2504  #ifdef TELINIT_USES_INITLVL
2505        FILE              *fp;
2506  #endif
2507        struct init_request     request;
```

init 2 于 init 3 的沟通通过“/dev/initctl”这个有名 pipe（其由 init 3 创建），在该 pipe 中传递的即是 init 2 对 init 3 提出的动作要求。即 init 2 本身不干任何事情，只是把用户的请求发送给 init 3 来处理。结构 init_request 就是“动作请求”的格式

```
struct init_request {
      int   magic;            /* Magic number                 */
      int   cmd;              /* What kind of request         */
      int   runlevel;         /* Runlevel to change to        */
      int   sleeptime;        /* Time between TERM and KILL   */
      union {
            struct init_request_bsd bsd;
            char              data[368];
      } i;
};
```

```
2508        struct sigaction  sa;
2509        int               f, fd, l;
2510        char              *env = NULL;
2511
2512        memset(&request, 0, sizeof(request));
2513        request.magic     = INIT_MAGIC;                 请求的签名
2514
```

下面就是根据命令行上的参数来构建 request

```
2515        while ((f = getopt(argc, argv, "t:e:")) != EOF) switch(f) {
2516              case 't':         t 参数指定 init 在向 process 发 SIGTERM 与 SIGKILL 之间的间隔秒数(sltime)
2517                    sltime = atoi(optarg);
2518                    break;
2519              case 'e':         手册中没有描述该参数，但看代码是用户可以通过该参数设置环境变量
```

```
2520                    if (env == NULL)
2521                          env = request.i.data;
2522                    l = strlen(optarg);
2523                    if (env + l + 2 > request.i.data + sizeof(request.i.data)) {
2524                          fprintf(stderr, "%s: -e option data "
2525                                "too large\n", progname);
2526                          exit(1);
2527                    }
2528                    memcpy(env, optarg, l);
2529                    env += l;
2530                    *env++ = 0;
2531                    break;
2532              default:
2533                    usage(progname);
2534                    break;
2535        }
2536
2537        if (env) *env++ = 0;
2538
2539        if (env) {
2540              if (argc != optind)
2541                    usage(progname);
2542              request.cmd = INIT_CMD_SETENV;
2543        } else {
2544              if (argc - optind != 1 || strlen(argv[optind]) != 1)
2545                    usage(progname);
2546              if (!strchr("0123456789SsQqAaBbCcUu", argv[optind][0]))
2547                    usage(progname);
2548              request.cmd = INIT_CMD_RUNLVL;            处理该 request 的代码在 check_init_fifo () 中(L2102)
2549              request.runlevel  = env ? 0 : argv[optind][0];  这里的 request.runlevel 即要切换过去的 run level
2550              request.sleeptime = sltime;
2551        }
2552
2553        /* Open the fifo and write a command. */
2554        /* Make sure we don't hang on opening /dev/initctl */
2555        SETSIG(sa, SIGALRM, signal_handler, 0);
2556        alarm(3);         设置 alarm，即 3 秒后本进程收到 SIGALRM signal
```

```
2557        if ((fd = open(INIT_FIFO, O_WRONLY)) >= 0 &&    打开有名管道“/dev/initctl”来写
2558            write(fd, &request, sizeof(request)) == sizeof(request)) {    写入该 request
2559              close(fd);
2560              alarm(0);          取消原先设置的 alarm 即上面 line 2556 行的 alarm(3)
                                     Line 2556 与 Line 2560 是为了保证对有名管道/dev/initctl 的写如果正常的话，肯定应该
2561              return 0;          小于 3 秒。如果超过 3 秒，则本进程会收到 SIGALRM signal，则在下面 Line 2585 中的
2562        }                        ISMEMBER（）将返回 true,则报错
2563
2564  #ifdef TELINIT_USES_INITLVL
2565        if (request.cmd == INIT_CMD_RUNLVL) {
2566              /* Fallthrough to the old method. */
2567
2568              /* Now write the new runlevel. */
2569              if ((fp = fopen(INITLVL, "w")) == NULL) {       打开/etc/initrunlvl 文件供写
2570                    fprintf(stderr, "%s: cannot create %s\n",
2571                          progname, INITLVL);
2572                    exit(1);
2573              }
2574              fprintf(fp, "%s %d", argv[optind], sltime);
2575              fclose(fp);
2576
2577              /* And tell init about the pending runlevel change. */
2578              if (kill(INITPID, SIGHUP) < 0) perror(progname);      向 1 号进程发 SIGHUP signal，也就是让
2579                                                                    init 2 重新读取/etc/inittab

2580              return 0;
2581        }
2582  #endif
2583
2584        fprintf(stderr, "%s: ", progname);
2585        if (ISMEMBER(got_signals, SIGALRM)) {
2586              fprintf(stderr, "timeout opening/writing control channel %s\n",
2587                    INIT_FIFO);
2588        } else {
2589              perror(INIT_FIFO);
2590        }
```

```
2591        return 1;
2592  }
```

下图可以表示 init 2 与 init 3 之间的关系。

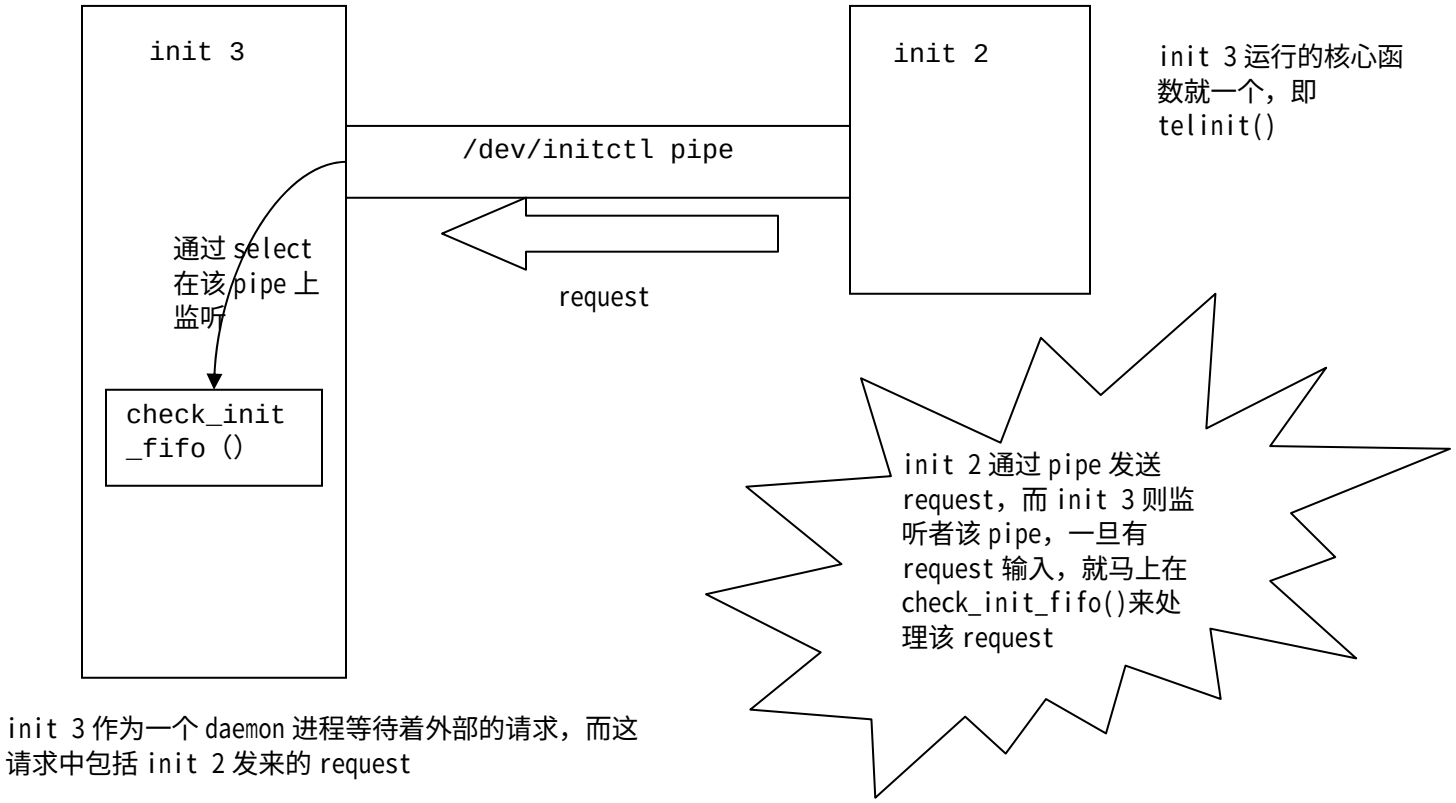

init 3 的运行

init 3 作为一个 daemon process 运行，它其实与 init 1 是完全一体的，即 init 1 完成系统用户态的启动，然后 init process 就进入类似 Client-Server 架构中的 server process 的运行。它本身不会主动的发起新的动作，而是当有 init process 关注的事件时，它才会被动的响应这些动作。我这里把原本连为一体的 init process 分成 init 1 与 init 3，纯粹是为了说明清楚的缘故。

主流程分析

当 init 1 结束后，即进入下面的一个无限循环。

```
2340  /*
2341   *    The main loop
2342   */
2343  int init_main()
2344  {

      。。。。。。    (这里省略的我把它归类为 init 1)

2456
2457    while(1) {      进入无限循环，init process 以后就在该无限循环中打转，永不出来，除非关机
2458
2459       /* See if we need to make the boot transitions. */
2460       boot_transitions();
2461       INITDBG(L_VB, "init_main: waiting..");
2462
2463       /* Check if there are processes to be waited on. */
2464       for(ch = family; ch; ch = ch->next)              枚举当前 init process 管理的进程
2465        if ((ch->flags & RUNNING) && ch->action != BOOT) break;     如果发觉有进程处于运行状态则马上退出枚举
```

```
2466
2467  #if CHANGE_WAIT
2468       /* Wait until we get hit by some signal. */
2469       while (ch != NULL && got_signals == 0) {
2470        if (ISMEMBER(got_signals, SIGHUP)) {
2471              /* See if there are processes to be waited on. */
2472              for(ch = family; ch; ch = ch->next)
2473                    if (ch->flags & WAITING) break;
2474        }
2475        if (ch != NULL) check_init_fifo();
2476       }
2477  #else /* CHANGE_WAIT */
2478       if (ch != NULL && got_signals == 0) check_init_fifo();
2479  #endif /* CHANGE_WAIT */
2480
2481       /* Check the 'failing' flags */
2482       fail_check();
2483
2484       /* Process any signals. */
2485       process_signals();
2486
2487       /* See what we need to start up (again) */
2488       start_if_needed();
2489     }
2490     /*NOTREACHED*/
2491  }
```

check_init_fifo () 的注解见“init 1”时的解释。这里是 init 3 检查是否有来自 init 2 的 request。

在 init 运行期间，其可能会收到关心的 signal。在 signal handler 中只是被纪录下来，而真正的处理则在该函数中。该函数注解见“init 1”的说明

该函数注解见“init 1”的说明

init process 等候着其关心的事件（通过 signal 的方式来通知），主要有如下几个事件

- 其通过 select 系统调用等候在管道**/dev/initctl** 的一端，一旦有进程往另一端写入希望 init process 执行的 request，init 3 就会分析该 request，然后执行要求的动作（见 check_init_fifo 函数注释）。
- 当有 signal 发送给 init process 时，init 在 signal handler 中并不马上处理该 signal，而仅仅是在全局变量 got_signals 中置一下该 signal 发生过的标志，真正的处理是在 init 3 的循环中的 process_signals()中。

设置 init process 的信号(signal)处理器。

```
  SETSIG(sa, SIGALRM,  signal_handler, 0);
```

```
  SETSIG(sa, SIGHUP,    signal_handler, 0);
  SETSIG(sa, SIGINT,    signal_handler, 0);
  SETSIG(sa, SIGPWR,    signal_handler, 0);
  SETSIG(sa, SIGWINCH, signal_handler, 0);
  SETSIG(sa, SIGUSR1,   signal_handler, 0);
```

在 init 1 中把 HUP PWR WINCH ALRM INT signal 的处理器设置好。当 signal 发生时，只是简单的置一下位，真正的对 signal 的处理是在函数 process_signals()中。

```
/*
 *    We got a signal (HUP PWR WINCH ALRM INT)
 */
void signal_handler(int sig)
{
      ADDSET(got_signals, sig);
}
```

而 ADDSET 是个 macro。简单设个标志位了事。

```
#define ADDSET(set, val)   ((set) |=  (1 << (val)))
```

在 process_signals()中检查哪些标志位置位，然后依次处理。

```
  if (ISMEMBER(got_signals, SIGPWR)) {          ISMEMBER 是个 macro，定义如下
      INITDBG(L_VB, "got SIGPWR");                    #define ISMEMBER(set, val) ((set) & (1 << (val)))
      /* See _what_ kind of SIGPWR this is. */  用于检查某个 signal 是否被触发
      pwrstat = 0;
      if ((fd = open(PWRSTAT, O_RDONLY)) >= 0) {
            c = 0;
            read(fd, &c, 1);                    电源 fail 的处理
            pwrstat = c;
            close(fd);
            unlink(PWRSTAT);
      }
      do_power_fail(pwrstat);
      DELSET(got_signals, SIGPWR);              处理完后把标志清调
  }

  if (ISMEMBER(got_signals, SIGINT)) {          当用户按了 Ctrl+Alt+Del 后，内核将发送 SIGINT signal 给 init
      INITDBG(L_VB, "got SIGINT");              process
```

```
    /* Tell ctrlaltdel entry to start up */
    for(ch = family; ch; ch = ch->next)
          if (ch->action == CTRLALTDEL)
                ch->flags &= ~XECUTED;        去掉 disable 执行的标志
    DELSET(got_signals, SIGINT);
}

if (ISMEMBER(got_signals, SIGWINCH)) {
    INITDBG(L_VB, "got SIGWINCH");
    /* Tell kbrequest entry to start up */
    for(ch = family; ch; ch = ch->next)
          if (ch->action == KBREQUEST)
                ch->flags &= ~XECUTED;
    DELSET(got_signals, SIGWINCH);
}

if (ISMEMBER(got_signals, SIGALRM)) {
    INITDBG(L_VB, "got SIGALRM");
    /* The timer went off: check it out */
    DELSET(got_signals, SIGALRM);
}

if (ISMEMBER(got_signals, SIGCHLD)) {
    INITDBG(L_VB, "got SIGCHLD");
    /* First set flag to 0 */
    DELSET(got_signals, SIGCHLD);

    /* See which child this was */
    for(ch = family; ch; ch = ch->next)
        if (ch->flags & ZOMBIE) {
          INITDBG(L_VB, "Child died, PID= %d", ch->pid);
          ch->flags &= ~(RUNNING|ZOMBIE|WAITING);
          if (ch->process[0] != '+')
                write_utmp_wtmp("", ch->id, ch->pid, DEAD_PROCESS, NULL);
        }

}
```

```
  if (ISMEMBER(got_signals, SIGHUP)) {
      INITDBG(L_VB, "got SIGHUP");
#if CHANGE_WAIT
      /* Are we waiting for a child? */
      for(ch = family; ch; ch = ch->next)
            if (ch->flags & WAITING) break;
      if (ch == NULL)
#endif
      {
            /* We need to go into a new runlevel */
            oldlevel = runlevel;
#ifdef INITLVL
            runlevel = read_level(0);
#endif
            if (runlevel == 'U') {
                  runlevel = oldlevel;
                  re_exec();
            } else {
                  if (oldlevel != 'S' && runlevel == 'S') console_stty();
                  if (runlevel == '6' || runlevel == '0' ||
                      runlevel == '1') console_stty();
                  read_inittab();
                  fail_cancel();
                  setproctitle("init [%c]", runlevel);
                  DELSET(got_signals, SIGHUP);
            }
      }
  }
  if (ISMEMBER(got_signals, SIGUSR1)) {
      /*
       *    SIGUSR1 means close and reopen /dev/initctl
       */
      INITDBG(L_VB, "got SIGUSR1");
      close(pipe_fd);
      pipe_fd = -1;
      DELSET(got_signals, SIGUSR1);
```

```
  }
```

■ 当受到 UPS 发来的电源 fail signal 时

```
      if ((fd = open(PWRSTAT, O_RDONLY)) >= 0) {
            c = 0;
            read(fd, &c, 1);
            pwrstat = c;
            close(fd);
            unlink(PWRSTAT);
      }
      do_power_fail(pwrstat);
```

首先从/etc/powerstatus 获得原因，“F”，“L”，“O”，分别表示电源 Fail，电源 Low，电源 Ok(恢复)。

参数 pwrstat 为 powerfail 原因，根据原因来启动/etc/inittab 中要求当对因事件发生时要执行的动作
参数 pwrstat 反映了电源（power）的状态
“O”表示电源 OK
“L”表示电源 Low
“F” 表示电源 Fail（故障）

```
1757  /*
1758   *    Start up powerfail entries.
1759   */
1760  void do_power_fail(int pwrstat)
1761  {
1762        CHILD *ch;
1763
1764        /*
1765         *    Tell powerwait & powerfail entries to start up
1766         */
1767        for (ch = family; ch; ch = ch->next) {    枚举 family 链表
1768              if (pwrstat == 'O') {               在收到电源恢复（Ok）的 signal 后
1769                    /*
1770                     *    The power is OK again.
1771                     */
```

```
1772                    if (ch->action == POWEROKWAIT)      XECUTED 是 disable 该 process 运行
1773                            ch->flags &= ~XECUTED;      清除 disable 标志，即属性为 POWEROKWAIT 的 process 可以
                                                            运行了
1774            } else if (pwrstat == 'L') {                在收到电源 Low 的 signal 后
1775                    /*
1776                     *      Low battery, shut down now.
1777                     */
1778                    if (ch->action == POWERFAILNOW)
1779                            ch->flags &= ~XECUTED;      允许属性为 POWERFAILNOW 的 process 的运行
1780            } else {
1781                    /*
1782                     *      Power is failing, shutdown imminent       在收到电源 Fail 的 signal 后
1783                     */
1784                    if (ch->action == POWERFAIL || ch->action == POWERWAIT)
1785                            ch->flags &= ~XECUTED;      允许属性为 POWERFAIL 与 POWERWAIT 的 process 的运行
1786            }
1787    }
1788  }
```

对应到例子文件中

```
# When our UPS tells us power has failed, assume we have a few minutes
# of power left.  Schedule a shutdown for 2 minutes from now.
# This does, of course, assume you have powerd installed and your
# UPS connected and working correctly.
pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"

# If power was restored before the shutdown kicked in, cancel it.
pr:12345:powerokwait:/sbin/shutdown -c "Power Restored; Shutdown Cancelled"
```

也就是在默认情况下/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"（属性为 POWERFAIL）与 /sbin/shutdown -c "Power Restored; Shutdown Cancelled"(属性为 POWEROKWAIT)都在 init process 管理的进程链表中是被 disable 的，即置了 XECUTED 标志的，当收到对应的消息后，会 enable 它。当然这里只是打开了允许他们运行，并没有真正的运行。真正的运行在下面的 start_if_needed()中。

```
2484        /* Process any signals. */
2485        process_signals();          处理被记录下来的 signal
```

```
2486
2487        /* See what we need to start up (again) */
2488        start_if_needed();          有可能在 family 链表中的 node 状态已经改变，所以重新枚举整个链表，看是否有本来不能运行
2489    }                               的动作(action)可以执行了
```

■ 当用户按了中断键（Delete 键或 Ctrl-C），则会向前台 process group 发 SIGINT signal。

```
2263        /* Tell ctrlaltdel entry to start up */
2264        for(ch = family; ch; ch = ch->next)
2265              if (ch->action == CTRLALTDEL)
2266                    ch->flags &= ~XECUTED;   允许 Ctrl-Alt-Del handler 的运行
```

■ 按《UNIX 环境高级编程》上的说法如下：
当有进程通过 ioctl 接口来改变终端窗口大小时会发该消息。

```
2272        /* Tell kbrequest entry to start up */
2273        for(ch = family; ch; ch = ch->next)
2274              if (ch->action == KBREQUEST)
2275                    ch->flags &= ~XECUTED;   允许属性为 KBREQUEST 的 process 运行
```

■ 对 timeout 的 alarm，几乎什么都不干

```
2280        INITDBG(L_VB, "got SIGALRM");
2281        /* The timer went off: check it out */
2282        DELSET(got_signals, SIGALRM);
```

■ 当收到有子进程死亡时，要在 utmp 和 wtmp 文件中记上一笔。

```
2290        /* See which child this was */
2291        for(ch = family; ch; ch = ch->next)       对 family 链表进行枚举，如果发现僵尸，则清除三个标志
2292            if (ch->flags & ZOMBIE) {
2293              INITDBG(L_VB, "Child died, PID= %d", ch->pid);
2294              ch->flags &= ~(RUNNING|ZOMBIE|WAITING);
2295              if (ch->process[0] != '+')
2296                    write_utmp_wtmp("", ch->id, ch->pid, DEAD_PROCESS, NULL);
2297            }
```

■ 当用户如果修改了/etc/inittab，希望不用重新启动就生效（在正常情况下，init process 只在启统启动阶段才会读取该配置文件），则可以发 SIGHUP signal，让 init process 重读 inittab，并根据新的配置文件来运行。

```
2303  #if CHANGE_WAIT
2304        /* Are we waiting for a child? */
2305        for(ch = family; ch; ch = ch->next)            如果这时有 process 处于必须等待其结束的状态，则此时不能
2306              if (ch->flags & WAITING) break;          执行重读 inittab 文件的操作
2307        if (ch == NULL)          如果有 process 处于必须等待其结束的状态，则 ch 就不会为 NULL，即如果为 NULL，表示此
2308  #endif                         时重新读取 inittab 文件是安全的
2309        {
2310              /* We need to go into a new runlevel */
2311              oldlevel = runlevel;          把当前的 run  level 存入 oldlevel 中
2312  #ifdef INITLVL
2313              runlevel = read_level(0);     获得当前 inittab 文件中的 run  level
2314  #endif
2315              if (runlevel == 'U') {        按照 init 手册的说法如下“U or u tell  init to re-execute itself
2316                    runlevel = oldlevel;    (preserving the state). No re-examining of /etc/inittab
2317                    re_exec();              file happens”，其中 re_exec()即时重新启动一遍 init 管理的链表中的
                                              Process
2318              } else {
2319                    if (oldlevel != 'S' && runlevel == 'S') console_stty();
2320                    if (runlevel == '6' || runlevel == '0' ||
2321                        runlevel == '1') console_stty();
2322                    read_inittab();         我们在分析 init  1 时说过，read_inittab()中有部分分支是不在系统时执行
2323                    fail_cancel();          的，是在 init  3 时执行的。
2324                    setproctitle("init [%c]", runlevel);
```

■ 当 init process 收到 SIGUSR1 signal(用户自定义信号)后，会关闭/dev/initctl pipe。

```
      close(pipe_fd);          init 在接受到该 signal 后，init 关闭和重新打开/dev/initctl
      pipe_fd = -1;
```

辅助函数

fail_check函数

```
1707  /*
1708   *    This procedure is called after every signal (SIGHUP, SIGALRM..)
1709   *
1710   *    Only clear the 'failing' flag if the process is sleeping
1711   *    longer than 5 minutes, or inittab was read again due
1712   *    to user interaction.
1713   */
1714  void fail_check(void)
1715  {
1716        CHILD *ch;                    /* Pointer to child structure */
1717        time_t      t;                /* System time */
1718        time_t      next_alarm = 0;   /* When to set next alarm */
1719
1720        time(&t);         取得当前系统时间
1721
1722        for(ch = family; ch; ch = ch->next) {     枚举 init process 管理的进程
1723
1724              if (ch->flags & FAILING) {          枚举整个 family 链表检查状态为 fail 的 node
1725                    /* Can we free this sucker? */
1726                    if (ch->tm + SLEEPTIME < t) {
1727                          ch->flags &= ~FAILING;
1728                          ch->count = 0;
1729                          ch->tm = 0;
1730                    } else {
1731                          /* No, we'll look again later */
1732                          if (next_alarm == 0 ||
1733                              ch->tm + SLEEPTIME > next_alarm)
1734                                next_alarm = ch->tm + SLEEPTIME;
1735                    }
1736              }
1737        }
```

```
1738        if (next_alarm) {
1739              next_alarm -= t;
1740              if (next_alarm < 1) next_alarm = 1;
1741              alarm(next_alarm);
1742        }
1743  }
```

re_exec函数

```
1827  /*
1828   *    Attempt to re-exec.
1829   */
1830  void re_exec(void)
1831  {
1832        CHILD       *ch;
1833        sigset_t    mask, oldset;
1834        pid_t       pid;
1835        char        **env;
1836        int         fd;
1837
1838        if (strchr("S12345",runlevel) == NULL)    这两行对应 telinit 手册中的说明“Run level should be one of
1839              return;                             Ss12345,otherwise request would be silently ignored“

1840
1841        /*
1842         *    Reset the alarm, and block all signals.
1843         */
1844        alarm(0);                       取消所有 alarm signal
1845        sigfillset(&mask);
1846        sigprocmask(SIG_BLOCK, &mask, &oldset);
1847
1848        /*
1849         *    construct a pipe fd --> STATE_PIPE and write a signature
1850         */
1851        fd = make_pipe(STATE_PIPE);
1852
```

```
1853        /*
1854         * It's a backup day today, so I'm pissed off.  Being a BOFH, however,
1855         * does have it's advantages...
1856         */
1857        fail_cancel();
1858        close(pipe_fd);
1859        pipe_fd = -1;
1860        DELSET(got_signals, SIGCHLD);
1861        DELSET(got_signals, SIGHUP);
1862        DELSET(got_signals, SIGUSR1);
1863
1864        /*
1865         *    That should be cleaned.
1866         */
1867        for(ch = family; ch; ch = ch->next)       对 init process 管理的进程链表进行枚举
1868            if (ch->flags & ZOMBIE) {
1869              INITDBG(L_VB, "Child died, PID= %d", ch->pid);
1870              ch->flags &= ~(RUNNING|ZOMBIE|WAITING);
1871              if (ch->process[0] != '+')
1872                    write_utmp_wtmp("", ch->id, ch->pid, DEAD_PROCESS, NULL);
1873            }
1874
1875        if ((pid = fork()) == 0) {
1876              /*
1877               *    Child sends state information to the parent.
1878               */
1879              send_state(fd);         由 init 的子进程往 state pipe 中写入当前 init process 管理的 process 的状态
1880              exit(0);
1881        }
1882
1883        /*
1884         *    The existing init process execs a new init binary.
1885         */
1886        env = init_buildenv(0);
1887        execl(myname, myname, "--init", NULL, env);     原有的 init process 执行/sbin/init，环境用
1888                                                        init_buildenv()构造的，而非父进程
```

```
1889        /*
1890         *    We shouldn't be here, something failed.
1891         *    Bitch, close the state pipe, unblock signals and return.
1892         */
1893        close(fd);
1894        close(STATE_PIPE);
1895        sigprocmask(SIG_SETMASK, &oldset, NULL);
1896        init_freeenv(env);
1897        initlog(L_CO, "Attempt to re-exec failed");
1898    }
```

read_inittab函数在init 3时的分析

```
1108  void read_inittab(void)
1109  {
1110    FILE            *fp;              /* The INITTAB file */
1111    CHILD           *ch, *old, *i;    /* Pointers to CHILD structure */
1112    CHILD           *head = NULL;     /* Head of linked list */
1113  #ifdef INITLVL
1114    struct stat     st;               /* To stat INITLVL */
1115  #endif
1116    sigset_t  nmask, omask;           /* For blocking SIGCHLD. */
1117    char            buf[256];         /* Line buffer */
1118    char            err[64];          /* Error message. */
1119    char            *id, *rlevel,
1120                *action, *process;    /* Fields of a line */
1121    char            *p;
1122    int       lineNo = 0;       /* Line number in INITTAB file */
1123    int       actionNo;         /* Decoded action field */
1124    int       f;                /* Counter */
1125    int       round;            /* round 0 for SIGTERM, 1 for SIGKILL */
1126    int       foundOne = 0;     /* No killing no sleep */
1127    int       talk;             /* Talk to the user */
1128    int       done = 0;         /* Ready yet? */
1129
1130  #if DEBUG
1131    if (newFamily != NULL) {
```

```
1132        INITDBG(L_VB, "PANIC newFamily != NULL");
1133        exit(1);
1134    }
1135    INITDBG(L_VB, "Reading inittab");
1136  #endif
1137
1138    /*
1139     *  Open INITTAB and real line by line.
1140     */
1141    if ((fp = fopen(INITTAB, "r")) == NULL)             打开/etc/inittab 文件
1142        initlog(L_VB, "No inittab file found");
1143
1144    while(!done) {              每循环一次即处理 inittab 中一行，构造 newFamily 链表。注意是 newFamily 链表，不是
1145        /*                      family 链表
1146         *    Add single user shell entry at the end.
1147         */
1148        if (fp == NULL || fgets(buf, sizeof(buf), fp) == NULL) {
1149              done = 1;         已经处理完毕，退出循环
1150              /*
1151               *    See if we have a single user entry.
1152               */
1153              for(old = newFamily; old; old = old->next)            ？？？
1154                    if (strpbrk(old->rlevel, "S")) break;
1155              if (old == NULL)
1156                    snprintf(buf, sizeof(buf), "~~:S:wait:%s\n", SULOGIN);
1157              else
1158                    continue;
1159        }
1160        lineNo++;
1161        /*
1162         *    Skip comments and empty lines
1163         */
1164        for(p = buf; *p == ' ' || *p == '\t'; p++)      忽略前导空白字符
1165              ;
1166        if (*p == '#' || *p == '\n') continue;          以“#”开头的为注释，忽略
1167
```

```
1168        /*
1169         *    Decode the fields
1170         */
```
分解 id:runlevels:action:process 中的 4 部分
```
1171        id =      strsep(&p, ":");    由于文件中的配置行的各部分用“：”分割，所以这里通过 strsep 来分别提取各部分内
1172        rlevel =  strsep(&p, ":");    容
1173        action =  strsep(&p, ":");
1174        process = strsep(&p, "\n");
1175
```
从下面的代码可以看到在 init manual 中没有标明的限制，比如命令行的长度不能太长，超过 127 个字符等
```
1176        /*
1177         *    Check if syntax is OK. Be very verbose here, to
1178         *    avoid newbie postings on comp.os.linux.setup :)
1179         */
1180        err[0] = 0;
1181        if (!id || !*id) strcpy(err, "missing id field");
1182        if (!rlevel)     strcpy(err, "missing runlevel field");
1183        if (!process)    strcpy(err, "missing process field");
1184        if (!action || !*action)
1185                    strcpy(err, "missing action field");
1186        if (id && strlen(id) > sizeof(utproto.ut_id))
1187              sprintf(err, "id field too long (max %d characters)",
1188                    (int)sizeof(utproto.ut_id));
1189        if (rlevel && strlen(rlevel) > 11)
1190              strcpy(err, "rlevel field too long (max 11 characters)");
1191        if (process && strlen(process) > 127)
1192              strcpy(err, "process field too long");
1193        if (action && strlen(action) > 32)
1194              strcpy(err, "action field too long");
1195        if (err[0] != 0) {
1196              initlog(L_VB, "%s[%d]: %s", INITTAB, lineNo, err);
1197              INITDBG(L_VB, "%s:%s:%s:%s", id, rlevel, action, process);
1198              continue;
1199        }
1200
1201        /*
```

```
1202          *    Decode the "action" field
1203          */
      init 允许的 action 类型都记录在 actions[]数组中，这里通过比较字符串来把其转换成数字型 identifier
1204        actionNo = -1;
1205        for(f = 0; actions[f].name; f++)
1206              if (strcasecmp(action, actions[f].name) == 0) {
1207                    actionNo = actions[f].act;
1208                    break;
1209              }
1210        if (actionNo == -1) {          碰到非法的 action(不在 actions[]数组中的)则忽略
1211              initlog(L_VB, "%s[%d]: %s: unknown action field",
1212                    INITTAB, lineNo, action);
1213              continue;
1214        }
1215
1216        /*
1217         *    See if the id field is unique
1218         */
```

配置行中的第一部分是所谓 identifier，必须唯一，但命名好像没有任何规定，可任意。已经处理过的配置行都被记录入 CHILD 的链表节点中，这里在处理当前行时检查一下已有节点中是否有与当前行的 id 相同的，如果有，则不是忽略该行，而是停止继续处理/etc/inittab 文件，可见 id 的唯一性是至关重要的

```
1219        for(old = newFamily; old; old = old->next) {
1220              if(strcmp(old->id, id) == 0 && strcmp(id, "~~")) {
1221                    initlog(L_VB, "%s[%d]: duplicate ID field \"%s\"",
1222                          INITTAB, lineNo, id);
1223                    break;
1224              }
1225        }
1226        if (old) continue;
1227
1228        /*
1229         *    Allocate a CHILD structure
1230         */
1231        ch = imalloc(sizeof(CHILD));        为当前配置行分配一个 CHILD node
```

```
1232
1233        /*
1234         *    And fill it in.          用当前配置行中的信息来填充 CHILD node
1235         */
1236        ch->action = actionNo;         action 类型
1237        strncpy(ch->id, id, sizeof(utproto.ut_id) + 1); /* Hack for different libs. */该行的唯一标示符
1238        strncpy(ch->process, process, sizeof(ch->process) - 1);      该行是要执行的命令行
1239        if (rlevel[0]) {               填上 run level
1240              for(f = 0; f < sizeof(rlevel) - 1 && rlevel[f]; f++) {
1241                    ch->rlevel[f] = rlevel[f];
1242                    if (ch->rlevel[f] == 's') ch->rlevel[f] = 'S';
1243              }
1244              strncpy(ch->rlevel, rlevel, sizeof(ch->rlevel) - 1);
1245        } else {     如果没有写 run level，则表示所有 run level 都要执行该行的 process 部分
1246              strcpy(ch->rlevel, "0123456789");
1247              if (ISPOWER(ch->action))
1248                    strcpy(ch->rlevel, "S0123456789");
1249        }
      下面是对 action 的处理
1250        /*
1251         *    We have the fake runlevel '#' for SYSINIT  and
1252         *    '*' for BOOT and BOOTWAIT.
1253         */

      从上面的注释看，SYSINIT action 用'#'表示，而 BOOT action 用'*'表示。而真正合法的 run level 是 0 到 9 加上'S'
      “ # ”与“ * ”表示在任何 run level 都要执行，另外 SYSINIT 的优先级是最高的，所以它应该比 BOOT 中的 action 先执行

1254        if (ch->action == SYSINIT) strcpy(ch->rlevel, "#");
1255        if (ch->action == BOOT || ch->action == BOOTWAIT)
1256              strcpy(ch->rlevel, "*");
1257
1258        /*
1259         *    Now add it to the linked list. Special for powerfail.
1260         */
```

在从/etc/inittab 中读取配置行并生成的链表的头为 newFamily。如果是系统的启动阶段，family 所表示的链表自然为空，如果只是通过运行 init 来切换 run level 等，则 family 记录的链表非空，也就是当前 init 通过上次读取/etc/inittab 后生成的链表。

```
1261        if (ISPOWER(ch->action)) {
1262
1263              /*
1264               *    Disable by default
1265               */
1266              ch->flags |= XECUTED;
1267
1268              /*
1269               *    Tricky: insert at the front of the list..
1270               */
1271              old = NULL;
1272              for(i = newFamily; i; i = i->next) {
1273                    if (!ISPOWER(i->action)) break;
1274                    old = i;
1275              }
1276              /*
1277               *    Now add after entry "old"
1278               */
1279              if (old) {
1280                    ch->next = i;
1281                    old->next = ch;
1282                    if (i == NULL) head = ch;
1283              } else {
1284                    ch->next = newFamily;
1285                    newFamily = ch;
```

如果 action 属于这几种（POWERWAIT，POWERFAIL，POWEROKWAIT，POWERFAILNOW，CTRLALTDEL），即与电源相关与用户按了 Ctrl+Alt+Del 键

设置不执行标志。在 startup()种如果检测到该配置行所代表的 action 的 flag 设置了 XECUTED，则忽略对该行的处理。这可以理解，因为符合 ISPOWER()的 action 都不是在正常情况下需要运行的，只有对应的事件确实发生了，才需要执行。比如，如果用户从来没有按下过 Ctrl+Alt+Del 键，自然根本不需要执行/etc/inittab 中 CTRLALTDEL Action 指定的 process。所以在默认情况下它是 disable 的(通过设置 XECUTED 标志)。当检测到按下 Ctrl+Alt+Del 键后，才需要 enable。
并且上述 action 被插在 family 链表的前面，这样如果它们需要执行的话，将先得到执行。可以理解，因为这几个动作都比较严重，所以优先权较高

```
1286                    if (ch->next == NULL) head = ch;
1287              }
1288        } else {  其他的 action 都插到尾部，KBREQUEST 默认是不执行的。从 init manual 上看 SIGWINCH  signal 会触发该动作
1289              /*
1290               *    Just add at end of the list
1291               */
1292              if (ch->action == KBREQUEST) ch->flags |= XECUTED;
1293              ch->next = NULL;
1294              if (head)
1295                    head->next = ch;
1296              else
1297                    newFamily = ch;
1298              head = ch;
1299        }
1300
1301        /*
1302         *    Walk through the old list comparing id fields
1303         */
1304        for(old = family; old; old = old->next)
1305              if (strcmp(old->id, ch->id) == 0) {
1306                    old->new = ch;
1307                    break;
1308              }
          到这里处理一行结束
1309    }
1310    /*
1311     *  We're done.
1312     */
1313    if (fp) fclose(fp);    关闭/etc/inittab 文件，对于 init 1，实际上基本上到此为止。
1314                           下面都是作为 daemon 进程的 init 即 init 3 的处理。我觉得整个代码应该整理得更清晰一点，由
                               kernel 启动 init(init 1)与作为 daemon 进程运行的 init（init 3）的逻辑应该分开，而不要像现
                               在一样，绞和在一块，比较乱

1315    /*
1316     *  Loop through the list of children, and see if they need to
1317     *  be killed.
```

```
1318     */
1319
1320    INITDBG(L_VB, "Checking for children to kill");
1321    for(round = 0; round < 2; round++) {
1322      talk = 1;
1323      for(ch = family; ch; ch = ch->next) {
1324        ch->flags &= ~KILLME;
1325
```

循环两遍，round 0 for SIGTERM, 1 for SIGKILL，即第一遍给要杀死的 process 发 SIGTERM signal，第二遍发 SIGKILL signal

由于是系统启动阶段运行 init（init 1），则这时的 family 链表为空，应该不进入循环，则这时 round = 0，talk = 1，foundOne = 0，代码应该跳转到 L1393 行执行。

在 init 3 中，family 的链表当然不为空，所以自然会进入循环

从 L1321 到 L1414 的代码是 init 3 执行的代码，init 1 不会执行。

这部分代码就是首先把当前 init 管理的 process(在 family 链表中)需要杀死的就 kill 掉，为重新 inittab 中要启动的 process 清理场地。

1. 对如下属性的 process 是不能杀死的，即忽略这些进程
   - BOOT
   - S --- Single User Mode
   - DEMAND
2. 杀死 process 的方式是这样的，先发送 SIGTERM signal，然后等待 sltime 秒（sltime 默认为 5 秒，但用户可以设定），最后再发 SIGKILL signal。

```
1326        /*
1327         *    Is this line deleted?
1328         */
1329        if (ch->new == NULL) ch->flags |= KILLME;
1330
1331        /*
1332         *    If the entry has changed, kill it anyway. Note that
1333         *    we do not check ch->process, only the "action" field.
1334         *    This way, you can turn an entry "off" immediately, but
1335         *    changes in the command line will only become effective
1336         *    after the running version has exited.
1337         */
1338        if (ch->new && ch->action != ch->new->action) ch->flags |= KILLME;
1339
1340        /*
```

```
1341          *    Only BOOT processes may live in all levels
1342          */
1343         if (ch->action != BOOT &&
1344             strchr(ch->rlevel, runlevel) == NULL) {
1345               /*
1346                *    Ondemand procedures live always,
1347                *    except in single user
1348                */
1349               if (runlevel == 'S' || !(ch->flags & DEMAND))
1350                     ch->flags |= KILLME;
1351         }
1352
1353         /*
1354          *    Now, if this process may live note so in the new list
1355          */
1356         if ((ch->flags & KILLME) == 0) {
1357               ch->new->flags  = ch->flags;
1358               ch->new->pid    = ch->pid;
1359               ch->new->exstat = ch->exstat;
1360               continue;
1361         }
1362
1363
1364         /*
1365          *    Is this process still around?
1366          */
1367         if ((ch->flags & RUNNING) == 0) {
1368               ch->flags &= ~KILLME;
1369               continue;
1370         }
1371         INITDBG(L_VB, "Killing \"%s\"", ch->process);
1372         switch(round) {                                    先发 SIGTERM signal
1373               case 0: /* Send TERM signal */
1374                     if (talk)
1375                           initlog(L_CO,
1376                                 "Sending processes the TERM signal");
1377                     kill(-(ch->pid), SIGTERM);          将 SIGTERM signal 发给其进程组 id 为 ch->pid 的
```

```
1378                    foundOne = 1;                          process
1379                    break;
1380              case 1: /* Send KILL signal and collect status */
1381                    if (talk)                              发送 SIGKILL signal
1382                          initlog(L_CO,
1383                                "Sending processes the KILL signal");
1384                    kill(-(ch->pid), SIGKILL);
1385                    break;
1386        }
1387        talk = 0;
1388
1389      }
1390      /*
1391       *      See if we have to wait 5 seconds
1392       */
1393      if (foundOne && round == 0) {    对于系统启动阶段运行的 init（init 1），round = 0,但 foundOne = 0，所以
1394        /*                             将不会进入该 if 分支，跳转到 L1419 执行
1395         *    Yup, but check every second if we still have children.
1396         */                              同样，init 3 将执行执行这里 L1397 到 L1411 的代码
1397        for(f = 0; f < sltime; f++) {
1398              for(ch = family; ch; ch = ch->next) {
1399                    if (!(ch->flags & KILLME)) continue;
1400                    if ((ch->flags & RUNNING) && !(ch->flags & ZOMBIE))
1401                          break;
1402              }
1403              if (ch == NULL) {
1404                    /*
1405                     *    No running children, skip SIGKILL
1406                     */
1407                    round = 1;
1408                    foundOne = 0; /* Skip the sleep below. */
1409                    break;
1410              }
1411              do_sleep(1);            整个循环将等待 1 * sltime 秒
1412        }
1413      }
```

```
1414    }
1415
1416    /*
1417     *  Now give all processes the chance to die and collect exit statuses.
1418     */
1419    if (foundOne) do_sleep(1);              init 1 运行时，foundOne = 0，所以不会睡眠一秒
1420    for(ch = family; ch; ch = ch->next)     这时的 family 链表为空，不进入该循环，跳转到 L1437 执行
1421        if (ch->flags & KILLME) {           init 3 将执行这里 L1420 到 L1430 的代码
1422              if (!(ch->flags & ZOMBIE))
1423                  initlog(L_CO, "Pid %d [id %s] seems to hang", ch->pid,
1424                          ch->id);
1425              else {
1426                  INITDBG(L_VB, "Updating utmp for pid %d [id %s]",
1427                          ch->pid, ch->id);
1428                  ch->flags &= ~RUNNING;
1429                  if (ch->process[0] != '+')
1430                    write_utmp_wtmp("", ch->id, ch->pid, DEAD_PROCESS, NULL);
1431              }
1432        }
1433
1434    /*
1435     *  Both rounds done; clean up the list.
1436     */
1437    sigemptyset(&nmask);
1438    sigaddset(&nmask, SIGCHLD);
1439    sigprocmask(SIG_BLOCK, &nmask, &omask);
1440    for(ch = family; ch; ch = old) {            init 1 运行时，family 链表还为空，所以不进入循环，跳转到 L1444
1441        old = ch->next;         init 3 运行时，需要先释放原来 family 链表中的节点，重读 inittab 文件后生成的链表由
                                    newFamily 指向
1442        free(ch);
1443    }
1444    family = newFamily;             newFamily 中就是在本函数头上读入的当前的/etc/inittab 配置行生成的链表，现在才
                                        把它赋给 family

1445    for(ch = family; ch; ch = ch->next) ch->new = NULL;
1446    newFamily = NULL;
```

```
1447    sigprocmask(SIG_SETMASK, &omask, NULL);
1448
1449  #ifdef INITLVL
1450    /*
1451     *  Dispose of INITLVL file.
1452     */ 删除/etc/initrunlvl，根据是文件还是 symbol link，删除方式是不一样的
1453    if (lstat(INITLVL, &st) >= 0 && S_ISLNK(st.st_mode)) {     检查/etc/initrunlvl 是文件还是 symbol link
1454        /*
1455         *    INITLVL is a symbolic link, so just truncate the file.
1456         */
1457        close(open(INITLVL, O_WRONLY|O_TRUNC));
1458    } else {
1459        /*
1460         *    Delete INITLVL file.
1461         */
1462        unlink(INITLVL);
1463    }
1464  #endif
1465  #ifdef INITLVL2
1466    /*
1467     *  Dispose of INITLVL2 file.
1468     */ 删除/var/log/initrunlvl，根据是文件还是 symbol link，删除方式是不一样的
1469    if (lstat(INITLVL2, &st) >= 0 && S_ISLNK(st.st_mode)) {
1470        /*
1471         *    INITLVL2 is a symbolic link, so just truncate the file.
1472         */
1473        close(open(INITLVL2, O_WRONLY|O_TRUNC));
1474    } else {
1475        /*
1476         *    Delete INITLVL2 file.
1477         */
1478        unlink(INITLVL2);
1479    }
1480  #endif
1481  }
```

后记

本笔记是我想写的Linux系统初始化系列的一部分。（主要分三部分，一是init进程分析，也就是本文；二是多用户模式启动分析，即进入console登录界面的启动；三是GUI模式启动，即X Window登录界面的启动[1]）Linux内核初始化后的动作是启动init process，而以后的用户态的初始化全部由其启动，所以它是研究系统用户态启动的源头，非学习不可。

[1] 这里不涉及Linux内核本身的初始化，也就是上面介绍的start_kernel函数及其调用的相关函数。因为这个话题实在太大，几乎涉及内核所有子系统。

联系

Walter Zhou

mailto:z-l-dragon@hotmail.com

附录

环境

1. sysvinit-2.86 包
2. VMware + Redhat 8.0

inittab 中 action 的注解

- **respawn**：如果process字段指定的进程不存在，则启动该进程，init不等待处理结束，而是继续扫描inittab文件中的后续进程，当这样的进程终止时，init会重新启动它，如果这样的进程已存在，则什么也不做。
- **wait**：启动process字段指定的进程，并等到处理结束才去处理inittab中的下一记录项。
- **once**：启动process字段指定的进程，不等待处理结束就去处理下一记录项。当这样的进程终止时，也不再重新启动它，在进入新的运行级别时，如果这样的进程仍在运行，init也不重新启动它。
- **boot**：只有在系统启动时，init才处理这样的记录项，启动相应进程，并不等待处理结束就去处理下一个记录项。当这样的进程终止时，系统也不重启它。
- **bootwait**：系统启动后，当第一次从单用户模式进入多用户模式时处理这样的记录项，init启动这样的进程，并且等待它的处理结束，然后再进行下一个记录项的处理，当这样的进程终止时，系统也不重启它。
- **powerfail**：当init接到断电的信号（SIGPWR）时，处理指定的进程。
- **powerwait**：当init接到断电的信号（SIGPWR）时，处理指定的进程，并且等到处理结束才去检查其他的记录项。
- **off**：如果指定的进程正在运行，init就给它发SIGTERM警告信号，在向它发出信号SIGKILL强制其结束之前等待 5 秒，如果这样的进程不存在，则忽略这一项。
- **ondemand**：功能通respawn，不同的是，与具体的运行级别无关，只用于rstate字段是a、b、c的那些记录项。
- **sysinit**：指定的进程在访问控制台之前执行，这样的记录项仅用于对某些设备的初始化，目的是为了使init在这样的设备上向用户提问有关运行级别的问题，init需要等待进程运行结束后才继续。

- **initdefault**：指定一个默认的运行级别，只有当init一开始被调用时才扫描这一项，如果rstate字段指定了多个运行级别，其中最大的数字是默认的运行级别，如果rstate字段是空的，init认为字段是 0123456，于是进入级别 6，这样便陷入了一个循环，如果inittab文件中没有包含initdefault的记录项，则在系统启动时请求用户为它指定一个初始运行级别

关机分析

在 init process 的配置文件 inittab 中有多个动作涉及到 shutdown(关机命令)。

```
# Trap CTRL-ALT-DELETE
ca::ctrlaltdel:/sbin/shutdown -t3 -r now                    当按下 CTRL-ALT-DELETE 的组合键

# When our UPS tells us power has failed, assume we have a few minutes
# of power left.  Schedule a shutdown for 2 minutes from now.
# This does, of course, assume you have powerd installed and your
# UPS connected and working correctly.                      当 UPS 报告电源 fail 时
pf::powerfail:/sbin/shutdown -f -h +2 "Power Failure; System Shutting Down"

# If power was restored before the shutdown kicked in, cancel it.       当 UPS 报告电源恢复
pr:12345:powerokwait:/sbin/shutdown -c "Power Restored; Shutdown Cancelled"
```

在上面三种情况下，init process 都会调用 shutdown 命令，只不过参数不同而已。比如当按下 CTRL-ALT-DELETE 的组合键后，“ -t3 -r now ”表示从现在(now)开始等待 3 秒(t3)，然后重启系统(r)。具体参数请看该命令的手册。

为什么在介绍系统启动的 init process 时要介绍关机命令呢？很简单，shutdown 命令与 init process 息息相关。shutdown 真正的关机或重启实际上是通过 init process 来实现的。

关机流程介绍

Shutdown 命令有如下的参数选项

```
/*
 *    Show usage message.
 */
void usage(void)
```

```
{
      fprintf(stderr,
      "Usage:\t  shutdown [-akrhHPfnc] [-t secs] time [warning message]\n"
      "\t\t  -a:      use /etc/shutdown.allow\n"
      "\t\t  -k:      don't really shutdown, only warn.\n"
      "\t\t  -r:      reboot after shutdown.\n"
      "\t\t  -h:      halt after shutdown.\n"
      "\t\t  -P:      halt action is to turn off power.\n"
      "\t\t  -H:      halt action is to just halt.\n"
      "\t\t  -f:      do a 'fast' reboot (skip fsck).\n"
      "\t\t  -F:      Force fsck on reboot.\n"
      "\t\t  -n:      do not go through \"init\" but go down real fast.\n"
      "\t\t  -c:      cancel a running shutdown.\n"
      "\t\t  -t secs: delay between warning and kill signal.\n"
      "\t\t  ** the \"time\" argument is mandatory! (try \"now\") **\n");
      exit(1);
}
```

-a 表示只有记录在/etc/shutdown.allow 文件中的用户才允许运行 shutdown 命令。
-k 表示并不是真正要关机或重启，仅仅发给登录该系统的用户警告。
-r 表示系统重新启动。
-h 表示系统 halt。
-P 表示系统关机。
-f 表示重新启动，同时重启时不运行 fsck(检查文件系统)
-F 表示重新启动，同时重启时强制运行 fsck(检查文件系统)
-c 表示取消已进入 shutdown 状态的系统。
-t secs: 指定在发警告与杀死系统中运行的进程之间的间隔秒数

与 shutdown 相关的一些配置文件。

/etc/nologin	如果存在该文件，则不允许登陆该系统。如果系统要禁止用户登录，则只要建立该文件即可。该文件只要存在及可，有无内容无所谓。

```
/*
 *    Create the /etc/nologin file.
 */
void donologin(int min)
```

```
{
      FILE *fp;
      time_t t;

      time(&t);
      t += 60 * min;

      if ((fp = fopen(NOLOGIN, "w")) != NULL) {        create /etc/nologin 文件
            fprintf(fp, "\rThe system is going down on %s\r\n", ctime(&t));
            if (message[0]) fputs(message, fp);
            fclose(fp);
      }
}
```

/fastboot	如果希望在重启时不要检查文件系统（fsck），就建立该文件，有无内容无所谓。

```
      while((c = getopt(argc, argv, "HPacqkrhnfFyt:g:i:")) != EOF) {     分析 shutdown 命令行参数
            switch(c) {
                  case 'H':
                        halttype = "HALT";
                        break;
                  case 'P':
                        halttype = "POWERDOWN";
                        break;
                  case 'a': /* Access control. */
                        useacl = 1;
                        break;
                  case 'c': /* Cancel an already running shutdown. */
                        cancel = 1;
                        break;
                  case 'k': /* Don't really shutdown, only warn.*/
                        dontshut = 1;
                        break;
                  case 'r': /* Automatic reboot */
                        down_level[0] = '6';
                        break;
```

```
                case 'h': /* Halt after shutdown */
                      down_level[0] = '0';
                      break;
                case 'f': /* Don't perform fsck after next boot */
                      fastboot = 1;      表示希望重启时不要检查文件系统，置标志
                      break;
```

```
      chdir("/");
      if (fastboot)  close(open(FASTBOOT,  O_CREAT | O_RDWR, 0644));     如果 fastboot 标志置位建立/fastboot 文件
```

/forcefsck	如果希望在重启时强制检查文件系统（fsck），就建立该文件，有无内容无所谓。

```
      if (forcefsck) close(open(FORCEFSCK, O_CREAT | O_RDWR, 0644));     如果 fastfsck 标志置位建立/forcefsck 文
                                                                          件
```

/etc/shutdown.allow	如果该文件存在，shutdown 命令会查看该文件，只有列在该文件中的用户才可以运行 shutdown 命令；如果没有该文件，则只有 root 用户才有权运行 shutdown 命令。

```
      /* Process the options. */
      while((c = getopt(argc, argv, "HPacqkrhnfFyt:g:i:")) != EOF) {
            switch(c) {
                case 'H':
                      halttype = "HALT";
                      break;
                case 'P':
                      halttype = "POWERDOWN";
                      break;
                case 'a': /* Access control. */      -a option，表示查看/etc/shutdown.allow 文件
                      useacl = 1;                    置标志
                      break;
      。。。。。。
```

```
      /* Do we need to use the shutdown.allow file ? */
      if (useacl && (fp = fopen(SDALLOW, "r")) != NULL) {    如果标志置位，则打开/etc/shutdown.allow 文件，里面每一
                                                              行是一个用户登录名
            /* Read /etc/shutdown.allow. */
            i = 0;
            while(fgets(buf, 128, fp)) {          循环读取一行
                  if (buf[0] == '#' || buf[0] == '\n') continue;   以#开头的为注释，忽略该行
                  if (i > 31) continue;           该文件最多可以有 32 个用户登录名，我不知道为什么要有这个限制
                  for(sp = buf; *sp; sp++) if (*sp == '\n') *sp = 0;     提取用户名
                  downusers[i++] = strdup(buf);        把提取出的用户名放入 downusers 数组
            }
            if (i < 32) downusers[i] = 0;
            fclose(fp);

            /* Now walk through /var/run/utmp to find logged in users. */
            while(!user_ok && (ut = getutent()) != NULL) {

                  /* See if this is a user process on a VC. */
                  if (ut->ut_type != USER_PROCESS) continue;
                  sprintf(term, "/dev/%.*s", UT_LINESIZE, ut->ut_line);
                  if (stat(term, &st) < 0) continue;
#ifdef major /* glibc */
                  if (major(st.st_rdev) != 4 ||
                      minor(st.st_rdev) > 63) continue;
#else
                  if ((st.st_rdev & 0xFFC0) != 0x0400) continue;
#endif
                  /* Root is always OK. */
                  if (strcmp(ut->ut_user, "root") == 0) {    root 用户总是可以关机的
                        user_ok++;
                        break;
                  }

                  /* See if this is an allowed user. */
                  for(i = 0; i < 32 && downusers[i]; i++)
                        if (!strncmp(downusers[i], ut->ut_user,    检查运行 shutdown 命令的用户名是否在 downusers
```

```
                        UT_NAMESIZE)) {                          数组中
                          user_ok++;        置位，可以运行 shutdown 命令标志
                          break;
                    }
            }
            endutent();

            /* See if user was allowed. */
            if (!user_ok) {
                  if ((fp = fopen(CONSOLE, "w")) != NULL) {
                        fprintf(fp, "\rshutdown: no authorized users "
                                    "logged in.\r\n");
                        fclose(fp);
                  }
                  exit(1);
            }
      }
```

Shutdown 源码

```
/*
 * shutdown.c     Shut the system down.
 *
 * Usage:   shutdown [-krhfnc] time [warning message]
 *            -k: don't really shutdown, only warn.
 *            -r: reboot after shutdown.
 *            -h: halt after shutdown.
 *            -f: do a 'fast' reboot (skip fsck).
 *            -F: Force fsck on reboot.
 *            -n: do not go through init but do it ourselves.
 *            -c: cancel an already running shutdown.
 *            -t secs: delay between SIGTERM and SIGKILL for init.
 *
 * Author:  Miquel van Smoorenburg, miquels@cistron.nl
```

```
 *
 * Version: @(#)shutdown  2.86-1  31-Jul-2004  miquels@cistron.nl
 *
 *          This file is part of the sysvinit suite,
 *          Copyright 1991-2004 Miquel van Smoorenburg.
 *
 *          This program is free software; you can redistribute it and/or
 *          modify it under the terms of the GNU General Public License
 *          as published by the Free Software Foundation; either version
 *          2 of the License, or (at your option) any later version.
 */
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/wait.h>
#include <time.h>
#include <string.h>
#include <unistd.h>
#include <errno.h>
#include <stdlib.h>
#include <stdio.h>
#include <signal.h>
#include <fcntl.h>
#include <stdarg.h>
#include <utmp.h>
#include <syslog.h>
#include "paths.h"
#include "reboot.h"
#include "initreq.h"

char *Version = "@(#) shutdown 2.86-1 31-Jul-2004 miquels@cistron.nl";

#define MESSAGELEN      256

int dontshut = 0; /* Don't shutdown, only warn  */
char down_level[2];     /* What runlevel to go to.    */
int dosync = 1;         /* Sync before reboot or halt */
int fastboot = 0; /* Do a 'fast' reboot         */
```

```
int forcefsck = 0;      /* Force fsck on reboot       */
char message[MESSAGELEN];     /* Warning message      */
char *sltime = 0; /* Sleep time                  */
char newstate[64];      /* What are we gonna do       */
int doself = 0;         /* Don't use init       */
int got_alrm = 0;

char *clean_env[] = {
      "HOME=/",
      "PATH=/bin:/usr/bin:/sbin:/usr/sbin",
      "TERM=dumb",
      NULL,
};

/* From "wall.c" */
extern void wall(char *, int, int);

/* From "utmp.c" */
extern void write_wtmp(char *user, char *id, int pid, int type, char *line);

/*
 *    Sleep without being interrupted.
 */
void hardsleep(int secs)
{
      struct timespec ts, rem;

      ts.tv_sec = secs;
      ts.tv_nsec = 0;

      while(nanosleep(&ts, &rem) < 0 && errno == EINTR)
            ts = rem;
}


/*
 *    Break off an already running shutdown.
 */
```

```
void stopit(int sig)
{
      unlink(NOLOGIN);
      unlink(FASTBOOT);
      unlink(FORCEFSCK);
      unlink(SDPID);
      printf("\r\nShutdown cancelled.\r\n");
      exit(0);
}

/*
 *    Show usage message.
 */
void usage(void)
{
      fprintf(stderr,
      "Usage:\t  shutdown [-akrhHPfnc] [-t secs] time [warning message]\n"
      "\t\t  -a:      use /etc/shutdown.allow\n"
      "\t\t  -k:      don't really shutdown, only warn.\n"
      "\t\t  -r:      reboot after shutdown.\n"
      "\t\t  -h:      halt after shutdown.\n"
      "\t\t  -P:      halt action is to turn off power.\n"
      "\t\t  -H:      halt action is to just halt.\n"
      "\t\t  -f:      do a 'fast' reboot (skip fsck).\n"
      "\t\t  -F:      Force fsck on reboot.\n"
      "\t\t  -n:      do not go through \"init\" but go down real fast.\n"
      "\t\t  -c:      cancel a running shutdown.\n"
      "\t\t  -t secs: delay between warning and kill signal.\n"
      "\t\t  ** the \"time\" argument is mandatory! (try \"now\") **\n");
      exit(1);
}


void alrm_handler(int sig)
{
      got_alrm = sig;
}
```

```
/*
 *    Set environment variables in the init process.
 */
int init_setenv(char *name, char *value)
{
      struct init_request     request;
      struct sigaction  sa;
      int               fd;
      int               nl, vl;

      memset(&request, 0, sizeof(request));
      request.magic = INIT_MAGIC;
      request.cmd = INIT_CMD_SETENV;
      nl = strlen(name);
      vl = value ? strlen(value) : 0;

      if (nl + vl + 3 >= sizeof(request.i.data))
            return -1;

      memcpy(request.i.data, name, nl);
      if (value) {
            request.i.data[nl] = '=';
            memcpy(request.i.data + nl + 1, value, vl);
      }

        /*
       *    Open the fifo and write the command.
         *  Make sure we don't hang on opening /dev/initctl
       */
      memset(&sa, 0, sizeof(sa));
      sa.sa_handler = alrm_handler;
      sigaction(SIGALRM, &sa, NULL);
      got_alrm = 0;
        alarm(3);
        if ((fd = open(INIT_FIFO, O_WRONLY)) >= 0 &&
```

```
            write(fd, &request, sizeof(request)) == sizeof(request)) {
                close(fd);
                alarm(0);
                return 0;
        }

        fprintf(stderr, "shutdown: ");
        if (got_alrm) {
                fprintf(stderr, "timeout opening/writing control channel %s\n",
                        INIT_FIFO);
        } else {
                perror(INIT_FIFO);
        }
        return -1;
}


/*
 *    Tell everyone the system is going down in 'mins' minutes.
 */
void warn(int mins)
{
      char buf[MESSAGELEN + sizeof(newstate)];
      int len;

      buf[0] = 0;
      strncat(buf, message, sizeof(buf) - 1);
      len = strlen(buf);

      if (mins == 0)
            snprintf(buf + len, sizeof(buf) - len,
                  "\rThe system is going down %s NOW!\r\n",
                  newstate);
      else
            snprintf(buf + len, sizeof(buf) - len,
                  "\rThe system is going DOWN %s in %d minute%s!\r\n",
                        newstate, mins, mins == 1 ? "" : "s");
```

```
	wall(buf, 1, 0);
}

/*
 *	Create the /etc/nologin file.
 */
void donologin(int min)
{
	FILE *fp;
	time_t t;

	time(&t);
	t += 60 * min;

	if ((fp = fopen(NOLOGIN, "w")) != NULL) {
		fprintf(fp, "\rThe system is going down on %s\r\n", ctime(&t));
		if (message[0]) fputs(message, fp);
		fclose(fp);
	}
}

/*
 *	Spawn an external program.
 */
int spawn(int noerr, char *prog, ...)
{
	va_list	ap;
	pid_t pid, rc;
	int	i;
	char	*argv[8];

	i = 0;
	while ((pid = fork()) < 0 && i < 10) {
		perror("fork");
		sleep(5);
		i++;
	}
```

```
      if (pid < 0) return -1;

      if (pid > 0) {
            while((rc = wait(&i)) != pid)
                  if (rc < 0 && errno == ECHILD)
                        break;
            return (rc == pid) ? WEXITSTATUS(i) : -1;
      }

      if (noerr) fclose(stderr);

      argv[0] = prog;
      va_start(ap, prog);
      for (i = 1; i < 7 && (argv[i] = va_arg(ap, char *)) != NULL; i++)
            ;
      argv[i] = NULL;
      va_end(ap);

      chdir("/");
      environ = clean_env;

      execvp(argv[0], argv);
      perror(argv[0]);
      exit(1);

      /*NOTREACHED*/
      return 0;
}

/*
 *    Kill all processes, call /etc/init.d/halt (if present)
 */
void fastdown()
{
      int do_halt = (down_level[0] == '0');
      int i;
```

```
#if 0
      char cmd[128];
      char *script;

      /*
       *    Currently, the halt script is either init.d/halt OR rc.d/rc.0,
       *    likewise for the reboot script. Test for the presence
       *    of either.
       */
      if (do_halt) {
            if (access(HALTSCRIPT1, X_OK) == 0)
                  script = HALTSCRIPT1;
            else
                  script = HALTSCRIPT2;
      } else {
            if (access(REBOOTSCRIPT1, X_OK) == 0)
                  script = REBOOTSCRIPT1;
            else
                  script = REBOOTSCRIPT2;
      }
#endif

      /* First close all files. */
      for(i = 0; i < 3; i++)
            if (!isatty(i)) {
                  close(i);
                  open("/dev/null", O_RDWR);
            }
      for(i = 3; i < 20; i++) close(i);
      close(255);

      /* First idle init. */
      if (kill(1, SIGTSTP) < 0) {
            fprintf(stderr, "shutdown: can't idle init.\r\n");
            exit(1);
      }
```

```
      /* Kill all processes. */
      fprintf(stderr, "shutdown: sending all processes the TERM signal...\r\n");
      kill(-1, SIGTERM);
      sleep(sltime ? atoi(sltime) : 3);
      fprintf(stderr, "shutdown: sending all processes the KILL signal.\r\n");
      (void) kill(-1, SIGKILL);

#if 0
      /* See if we can run /etc/init.d/halt */
      if (access(script, X_OK) == 0) {
            spawn(1, cmd, "fast", NULL);
            fprintf(stderr, "shutdown: %s returned - falling back "
                        "on default routines\r\n", script);
      }
#endif

      /* script failed or not present: do it ourself. */
      sleep(1); /* Give init the chance to collect zombies. */

      /* Record the fact that we're going down */
      write_wtmp("shutdown", "~~", 0, RUN_LVL, "~~");

      /* This is for those who have quota installed. */
      spawn(1, "accton", NULL);
      spawn(1, "quotaoff", "-a", NULL);

      sync();
      fprintf(stderr, "shutdown: turning off swap\r\n");
      spawn(0, "swapoff", "-a", NULL);
      fprintf(stderr, "shutdown: unmounting all file systems\r\n");
      spawn(0, "umount", "-a", NULL);

      /* We're done, halt or reboot now. */
      if (do_halt) {
            fprintf(stderr, "The system is halted. Press CTRL-ALT-DEL "
                        "or turn off power\r\n");
            init_reboot(BMAGIC_HALT);
```

```
            exit(0);
      }

      fprintf(stderr, "Please stand by while rebooting the system.\r\n");
      init_reboot(BMAGIC_REBOOT);
      exit(0);
}

/*
 *    Go to runlevel 0, 1 or 6.
 */
void shutdown(char *halttype)
{
      char  *args[8];
      int   argp = 0;
      int   do_halt = (down_level[0] == '0');

      /* Warn for the last time */
      warn(0);
      if (dontshut) {
            hardsleep(1);
            stopit(0);
      }
      openlog("shutdown", LOG_PID, LOG_USER);
      if (do_halt)
            syslog(LOG_NOTICE, "shutting down for system halt");
      else
            syslog(LOG_NOTICE, "shutting down for system reboot");
      closelog();

      /* See if we have to do it ourself. */
      if (doself) fastdown();

      /* Create the arguments for init. */
      args[argp++] = INIT;
      if (sltime) {
            args[argp++] = "-t";
```

```
            args[argp++] = sltime;
      }
      args[argp++] = down_level;
      args[argp]   = (char *)NULL;

      unlink(SDPID);
      unlink(NOLOGIN);

      /* Now execute init to change runlevel. */
      sync();
      init_setenv("INIT_HALT", halttype);
      execv(INIT, args);

      /* Oops - failed. */
      fprintf(stderr, "\rshutdown: cannot execute %s\r\n", INIT);
      unlink(FASTBOOT);
      unlink(FORCEFSCK);
      init_setenv("INIT_HALT", NULL);
      openlog("shutdown", LOG_PID, LOG_USER);
      syslog(LOG_NOTICE, "shutdown failed");
      closelog();
      exit(1);
}

/*
 *    returns if a warning is to be sent for wt
 */
static int needwarning(int wt)
{
      int ret;

      if (wt < 10)
            ret = 1;
      else if (wt < 60)
            ret = (wt % 15 == 0);
      else if (wt < 180)
            ret = (wt % 30 == 0);
```

```
      else
            ret = (wt % 60 == 0);

      return ret;
}

/*
 *    Main program.
 *    Process the options and do the final countdown.
 */
int main(int argc, char **argv)
{
      FILE              *fp;
      extern int        getopt();
      extern int        optind;
      struct sigaction  sa;
      struct tm         *lt;
      struct stat       st;
      struct utmp       *ut;
      time_t                  t;
      uid_t             realuid;
      char              *halttype;
      char              *downusers[32];
      char              buf[128];
      char              term[UT_LINESIZE + 6];
      char              *sp;
      char              *when = NULL;
      int               c, i, wt;
      int               hours, mins;
      int               didnolog = 0;
      int               cancel = 0;
      int               useacl = 0;
      int               pid = 0;
      int               user_ok = 0;

      /* We can be installed setuid root (executable for a special group) */
      realuid = getuid();
```

```
	setuid(geteuid());

	if (getuid() != 0) {
		fprintf(stderr, "shutdown: you must be root to do that!\n");
		exit(1);
	}
	strcpy(down_level, "1");
	halttype = NULL;

	/* Process the options. */
	while((c = getopt(argc, argv, "HPacqkrhnfFyt:g:i:")) != EOF) {
		switch(c) {
			case 'H':
				halttype = "HALT";
				break;
			case 'P':
				halttype = "POWERDOWN";
				break;
			case 'a': /* Access control. */
				useacl = 1;
				break;
			case 'c': /* Cancel an already running shutdown. */
				cancel = 1;
				break;
			case 'k': /* Don't really shutdown, only warn.*/
				dontshut = 1;
				break;
			case 'r': /* Automatic reboot */
				down_level[0] = '6';
				break;
			case 'h': /* Halt after shutdown */
				down_level[0] = '0';
				break;
			case 'f': /* Don't perform fsck after next boot */
				fastboot = 1;
				break;
			case 'F': /* Force fsck after next boot */
```

```
                        forcefsck = 1;
                        break;
                case 'n': /* Don't switch runlevels. */
                        doself = 1;
                        break;
                case 't': /* Delay between TERM and KILL */
                        sltime = optarg;
                        break;
                case 'y': /* Ignored for sysV compatibility */
                        break;
                case 'g': /* sysv style to specify time. */
                        when = optarg;
                        break;
                case 'i': /* Level to go to. */
                        if (!strchr("0156aAbBcCsS", optarg[0])) {
                                fprintf(stderr,
                                "shutdown: `%s': bad runlevel\n",
                                optarg);
                                exit(1);
                        }
                        down_level[0] = optarg[0];
                        break;
                default:
                        usage();
                        break;
        }
}

/* Do we need to use the shutdown.allow file ? */
if (useacl && (fp = fopen(SDALLOW, "r")) != NULL) {

        /* Read /etc/shutdown.allow. */
        i = 0;
        while(fgets(buf, 128, fp)) {
                if (buf[0] == '#' || buf[0] == '\n') continue;
                if (i > 31) continue;
                for(sp = buf; *sp; sp++) if (*sp == '\n') *sp = 0;
```

```
                downusers[i++] = strdup(buf);
            }
            if (i < 32) downusers[i] = 0;
            fclose(fp);

            /* Now walk through /var/run/utmp to find logged in users. */
            while(!user_ok && (ut = getutent()) != NULL) {

                  /* See if this is a user process on a VC. */
                  if (ut->ut_type != USER_PROCESS) continue;
                  sprintf(term, "/dev/%.*s", UT_LINESIZE, ut->ut_line);
                  if (stat(term, &st) < 0) continue;
#ifdef major /* glibc */
                  if (major(st.st_rdev) != 4 ||
                      minor(st.st_rdev) > 63) continue;
#else
                  if ((st.st_rdev & 0xFFC0) != 0x0400) continue;
#endif
                  /* Root is always OK. */
                  if (strcmp(ut->ut_user, "root") == 0) {
                        user_ok++;
                        break;
                  }

                  /* See if this is an allowed user. */
                  for(i = 0; i < 32 && downusers[i]; i++)
                        if (!strncmp(downusers[i], ut->ut_user,
                            UT_NAMESIZE)) {
                              user_ok++;
                              break;
                        }
            }
            endutent();

            /* See if user was allowed. */
            if (!user_ok) {
                  if ((fp = fopen(CONSOLE, "w")) != NULL) {
```

```
                        fprintf(fp, "\rshutdown: no authorized users "
                                    "logged in.\r\n");
                        fclose(fp);
                  }
                  exit(1);
            }
      }

      /* Read pid of running shutdown from a file */
      if ((fp = fopen(SDPID, "r")) != NULL) {
            fscanf(fp, "%d", &pid);
            fclose(fp);
      }

      /* Read remaining words, skip time if needed. */
      message[0] = 0;
      for(c = optind + (!cancel && !when); c < argc; c++) {
            if (strlen(message) + strlen(argv[c]) + 4 > MESSAGELEN)
                  break;
            strcat(message, argv[c]);
            strcat(message, " ");
      }
      if (message[0]) strcat(message, "\r\n");

      /* See if we want to run or cancel. */
      if (cancel) {
            if (pid <= 0) {
                  fprintf(stderr, "shutdown: cannot find pid "
                              "of running shutdown.\n");
                  exit(1);
            }
            init_setenv("INIT_HALT", NULL);
            if (kill(pid, SIGINT) < 0) {
                  fprintf(stderr, "shutdown: not running.\n");
                  exit(1);
            }
            if (message[0]) wall(message, 1, 0);
```

```
            exit(0);
      }

      /* Check syntax. */
      if (when == NULL) {
            if (optind == argc) usage();
            when = argv[optind++];
      }

      /* See if we are already running. */
      if (pid > 0 && kill(pid, 0) == 0) {
            fprintf(stderr, "\rshutdown: already running.\r\n");
            exit(1);
      }

      /* Extra check. */
      if (doself && down_level[0] != '0' && down_level[0] != '6') {
            fprintf(stderr,
            "shutdown: can use \"-n\" for halt or reboot only.\r\n");
            exit(1);
      }

      /* Tell users what we're gonna do. */
      switch(down_level[0]) {
            case '0':
                  strcpy(newstate, "for system halt");
                  break;
            case '6':
                  strcpy(newstate, "for reboot");
                  break;
            case '1':
                  strcpy(newstate, "to maintenance mode");
                  break;
            default:
                  sprintf(newstate, "to runlevel %s", down_level);
                  break;
      }
```

```
	/* Create a new PID file. */
	unlink(SDPID);
	umask(022);
	if ((fp = fopen(SDPID, "w")) != NULL) {
		fprintf(fp, "%d\n", getpid());
		fclose(fp);
	} else if (errno != EROFS)
		fprintf(stderr, "shutdown: warning: cannot open %s\n", SDPID);

	/*
	 *	Catch some common signals.
	 */
	signal(SIGQUIT, SIG_IGN);
	signal(SIGCHLD, SIG_IGN);
	signal(SIGHUP,  SIG_IGN);
	signal(SIGTSTP, SIG_IGN);
	signal(SIGTTIN, SIG_IGN);
	signal(SIGTTOU, SIG_IGN);

	memset(&sa, 0, sizeof(sa));
	sa.sa_handler = stopit;
	sigaction(SIGINT, &sa, NULL);

	/* Go to the root directory */
	chdir("/");
	if (fastboot)  close(open(FASTBOOT,  O_CREAT | O_RDWR, 0644));
	if (forcefsck) close(open(FORCEFSCK, O_CREAT | O_RDWR, 0644));

	/* Alias now and take care of old '+mins' notation. */
	if (!strcmp(when, "now")) strcpy(when, "0");
	if (when[0] == '+') when++;

	/* Decode shutdown time. */
	for (sp = when; *sp; sp++) {
		if (*sp != ':' && (*sp < '0' || *sp > '9'))
			usage();
```

```
	}
	if (strchr(when, ':') == NULL) {
		/* Time in minutes. */
		wt = atoi(when);
		if (wt == 0 && when[0] != '0') usage();
	} else {
		/* Time in hh:mm format. */
		if (sscanf(when, "%d:%2d", &hours, &mins) != 2) usage();
		if (hours > 23 || mins > 59) usage();
		time(&t);
		lt = localtime(&t);
		wt = (60*hours + mins) - (60*lt->tm_hour + lt->tm_min);
		if (wt < 0) wt += 1440;
	}
	/* Shutdown NOW if time == 0 */
	if (wt == 0) shutdown(halttype);

	/* Give warnings on regular intervals and finally shutdown. */
	if (wt < 15 && !needwarning(wt)) warn(wt);
	while(wt) {
		if (wt <= 5 && !didnolog) {
			donologin(wt);
			didnolog++;
		}
		if (needwarning(wt)) warn(wt);
		hardsleep(60);
		wt--;
	}
	shutdown(halttype);

	return 0; /* Never happens */
}
```